

Angular Sparsity Is an Atomic-Sector Property: The Abelian Residue of Multi-Center Bonding*

J. Loutey¹

¹*Independent Researcher, Kent, Washington*

(Dated: August 10, 2026)

The GEOVAC program builds molecular Hamiltonians from atom-centered Coulomb-Sturmian blocks whose angular matrix elements are generated from quantum-number labels by Gaunt selection, with no spatial quadrature. This yields structurally sparse qubit Hamiltonians—334 Pauli terms for LiH at $Q = 30$, $O(Q^{2.5})$ scaling (Paper 14)—at the cost of a documented binding failure for second-row diatomics (the W1e wall, Papers 19/20).

We settle the question of whether that trade was necessary. Using exact rational two-center machinery (Mulliken–Ruedenberg auxiliary functions for overlaps; a hydrogenic eigen-trick for cross-center one-body terms, certified by an exact bra/ket identity holding entrywise), we census the *genuine* bare tensor (S, h, g) for a real fragment pair and find: **[MEASURED]** the Gaunt l -selection rule does not survive cross-center evaluation. Only the axial m -rule does. Electron-repulsion density inflates $13.8\times$ over the builder at $n_{\max} = 2$, rising to $15.0\times$ at $n_{\max} = 3$, and cross-center one-body entries measure $0.19\text{--}0.80$ Ha—bond scale, not perturbative.

We give the mechanism **[INTERNAL THEOREM]**: m labels the *abelian* symmetry shared by both centers (rotation about the internuclear axis is one global one-parameter group, so m is a common label), while l labels irreducible representations of a rotation group attached to a single center. The intersection of the two centers’ rotation groups is exactly the axial $U(1)$, whose character group is $\mathbb{Z}$ by Pontryagin duality. Abelian selection rules are additive charge conservation and survive; non-abelian Clebsch–Gordan selection requires the full sphere and does not. Surviving label-generated sparsity is therefore governed by the largest common abelian symmetry, which is computable from the molecular point group in advance.

Two re-readings follow. First, the $17.9\times$ Pauli inflation previously attributed to Löwdin orthogonalization (v4.73.1) was never orthogonalization’s price: it is within measurement of the genuine tensor’s own density, so Löwdin was paying an honest bill. Second **[OBSERVATION]**, four independent attempts to relocate the non-orthogonality—into the operator, into a dual basis, into a high-symmetry embedding, into device-measured many-body overlaps—each moved the cost without reducing it, because the cost measures the physical asymmetry of a two-center problem.

We close with a demonstration **[PANEL-VERIFIED]**: all-electron 12-electron NaH, the framework’s hardest documented failure, binds at $R_{\text{eq}} = 3.736 a_0$ ($+4.8\%$ vs experiment) with 91.1% of in-basis FCI binding from three fragment-native determinants, once the integrals are genuine end-to-end. This localizes W1e to the Hamiltonian specification rather than the correlation treatment, by construction.

I. INTRODUCTION: WHAT THE SPARSITY BOUGHT, AND WHAT IT COST

The GEOVAC framework’s computational proposition is structural sparsity. Angular matrix elements are generated from the quantum-number labels (n, l, m) by Gaunt selection rules and Wigner $3j$ symbols with no spatial integration, so the resulting second-quantized Hamiltonian is sparse by construction rather than by numerical thresholding. The downstream consequence is the qubit-resource result of Paper 14 [2]: composed LiH at 334 Pauli terms on $Q = 30$ qubits, $O(Q^{2.5})$ Pauli scaling against $O(Q^{3.9-4.3})$ for Gaussian baselines, and a $190\times$ reduction against cc-pVDZ.

The framework’s documented failure is chemical binding for second-row diatomics. Named W1e, it has resisted an unusually thorough sequence of attacks: six Phillips–Kleinman modifications, basis-level Schmidt orthogo-

nalization, core-correlation estimates, basis enlargement to $n_{\max} = 4$, an explicit-core Hartree–Fock treatment with twelve electrons, and a DMRG cross-check that reproduced FCI bit-identically at finite sector dimension without producing an interior minimum (Papers 19/20; CHANGELOG v2.19.4–v3.87.0). The wall was localized to the Hamiltonian specification—the projection from the continuous Level-4 solver to second-quantized $(h_1, \text{eri}, e_{\text{core}})$ integrals—rather than to the treatment of correlation.

This paper answers the question that localization leaves open. The framework neglects the overlap between orbitals on different centers, treating a non-orthogonal basis as though it were orthonormal, and that neglect is what preserves the sparsity. Was there an alternative? Could one keep the genuine cross-center physics *and* the label-generated sparsity, by declining to transform the integral tensor at all?

The answer is no, and the reason is group-theoretic rather than technical.

* Paper 58 in the Geometric Vacuum series

II. THE EXACT-ARITHMETIC CENSUS

A. Machinery and certificate

Two-center overlaps at arbitrary (n, l, m, Z) are computed as exact rationals via the Mulliken–Ruedenberg auxiliary functions A_n/B_n [3, 4], evaluated in prolate spheroidal coordinates where the two-center problem separates [1]. Because the result is rational, vanishing is *decidable* rather than inferred from a numerical tolerance—a property no quadrature scheme can offer, and the reason this census can make a claim about exact zeros at all.

Cross-center one-body elements are obtained by extending the same route with $1/r_A$ and $1/r_B$ kernels, each of which cancels against the volume element $(\xi^2 - \eta^2)$, followed by the hydrogenic eigen-trick. The assembly carries its own certificate: the bra/ket identity

$$E_a S - Z_B I_B = E_b S - Z_A I_A \quad (1)$$

holds as an *exact rational equality* on every entry, which is a machinery check no fortuitous cancellation survives. Four independent quadrature cross-checks agree at $\sim 2 \times 10^{-16}$.

B. Result

The census configuration is $Z_A = 3$ at the origin, $Z_B = 1$ at $R = 3\hat{z}$, hydrogenic orbitals to $n_{\max} = 2$ on each center, $M = 10$ spatial orbitals. Table I compares the genuine bare tensor against the production builder’s.

The electron-repulsion tensor is 29.4% dense genuinely against 2.1% for the builder. At $n_{\max} = 3$ ($M = 28$) the comparison is 114,280 nonzeros (18.6%) against 7,600 (1.2%): an inflation factor of $13.8\times$ rising to $15.0\times$, i.e. *increasing* with basis size.

Two further measurements matter for interpretation. First, per-side Gaunt feasibility removes less than 6% of entries beyond what the m -rule already removes, and the three cross classes $(AA|AB) + (AB|AB) + (AB|BB)$ account for 80% of the genuine tensor. The residual l -selection is therefore confined to the same-center $(AA|AA)$ and $(BB|BB)$ blocks. Second, cross-center magnitudes are large: overlaps 0.08–0.52, one-body entries 0.19–0.80 Ha. For scale, the entire NaH bond is

TABLE I. Genuine versus builder sparsity, $Z_A = 3/Z_B = 1$ at $R = 3a_0$, $n_{\max} = 2$ per center ($M = 10$). “Genuine zeros are” records which selection rule accounts for the surviving zeros.

object	genuine	builder	dense	genuine zeros are
S	32	10	100	m -rule only
h	44	22	100	m -rule only
g	2,944	214	10,000	m -rule + same-center Gaunt

under 2 eV, so individual entries in the dense part of the tensor exceed the binding energy being computed. They are not droppable as small corrections—an independent quantification of the W1d/W1e diagnosis.

III. MECHANISM: THE ABELIAN RESIDUE

The census result is not an accident of the chosen configuration. It is forced, and the forcing is elementary once stated in the right language.

Theorem 1 (Angular selection survives molecularization only in the abelian sector). *Let two nuclei A, B lie on the z -axis, with atom-centered basis functions labeled (n, l, m) about their own center, and let $\hat{O}$ be invariant under the molecular symmetry group. Then:*

- (i) *The m -selection rule holds exactly across centers: $\langle ab|\hat{O}|cd\rangle = 0$ unless $m_a + m_b = m_c + m_d$.*
- (ii) *No l -selection rule holds between basis functions on different centers.*
- (iii) *The surviving label-generated sparsity is exactly that generated by the largest symmetry group common to both centers, which is the axial $U(1)$ and is abelian.*

Proof sketch. For (i): rotation about the internuclear axis is a *single* one-parameter group acting on all of $\mathbb{R}^3$. Its generator L_z is independent of which point on the axis is taken as origin, so m is a global label shared by both centers rather than a per-center one. Invariance of $\hat{O}$ under this group gives additive conservation of m .

For (ii): l labels irreducible representations of the rotation group about a *specific* center. Rotations about A and rotations about B are distinct subgroups of the Euclidean group whose intersection is precisely the axial subgroup. A function that is an l -eigenfunction about B , re-expanded about A , has support on all l : the two-center re-expansion does not terminate. Hence no l -selection between cross-center pairs. (The same-center blocks retain Gaunt selection, consistent with the $< 6\%$ residual measured in Sec. II.)

For (iii): the intersection in (ii) is $SO(2) \cong U(1)$, abelian. Abelian compact groups have one-dimensional irreducibles labeled by characters, and by Pontryagin duality the character group of $U(1)$ is $\mathbb{Z}$ —which is the integer label m . Selection rules from abelian symmetry are additive conservation laws on integer charges, and are preserved by any operation respecting the axis. Non-abelian selection rules are Clebsch–Gordan couplings requiring the full group; reducing the group destroys them. $\square$

Theorem 1 converts a measurement into a prediction, because the largest common abelian symmetry is readable from the molecular point group before any integral is computed.

Corollary 1 (Sparsity budget from the point group). *For a molecule with point group G , label-generated angular sparsity is bounded by the grading available from the abelian structure of G . Linear molecules ($C_{\infty v}$, $D_{\infty h}$) carry a continuous $U(1)$ and therefore a $\mathbb{Z}$ grading, worth a multiplicative factor. Bent molecules carry only a finite point group and therefore a finite grading.*

Prediction 1 (Water, untested). *H_2O has point group C_{2v} , which is order four and abelian, isomorphic to $\mathbb{Z}_2 \times \mathbb{Z}_2$. All its irreducibles are one-dimensional, so it admits a clean two-bit grading and exactly two qubits of symmetry-based reduction—not a multiplicative factor. This sharpens the qualitative expectation that the framework’s angular advantage degrades for non-linear molecules into a specific count, and is falsifiable on the existing H_2O build.*

The measured factor of 3–5 for the m -rule at the bases of Sec. II is the $\mathbb{Z}$ -grading case of Corollary 1. It is exact and structural, and it is a constant factor—not the $O(Q^{2.5})$ scaling story the framework’s qubit results rest on.

IV. TWO RE-READINGS

A. Löwdin was paying an honest bill

The corpus previously recorded a $17.9\times$ Pauli inflation from retrofitting Löwdin orthogonalization onto the balanced LiH integrals (CHANGELOG v4.73.1), and a $14\times$ inflation in the earlier heterogeneous-nested experiment. Both were read as the price of orthogonalization, and orthogonalization was accordingly treated as the sparsity-destroying move.

The census shows this reading was wrong in an instructive way. The genuine tensor’s own inflation over the builder is $\sim 15\times$, within measurement of the Löwdin figure. Orthogonalization was not destroying sparsity that the untransformed tensor possessed; the untransformed tensor never possessed it. “Löwdin destroys sparsity” and “the molecular tensor is intrinsically l -dense cross-center” are the same fact observed twice, from two sides.

This matters because it removes an entire class of proposed remedies. Any scheme whose premise is “avoid the transformation and keep the sparsity” is answering a question that was mis-posed.

B. Cost conservation

Observation 1 (Relocating the metric does not reduce it). *Four independent attempts to move the non-orthogonality out of the way each succeeded in moving it and failed to reduce it:*

1. into the **operator** (Löwdin $S^{-1/2}$): operator densifies $17.9\times$;

2. into a **dual basis** (biorthogonal encoding): densifies comparably, and by theorem rather than happenstance [5];
3. into a **high-symmetry embedding** (hyperspherical Level $4N$): angular basis requirement explodes—LiH at 63.5% R_{eq} error with $l_{max} = 2$ genuinely insufficient;
4. into **device-measured many-body overlaps** (NOCI): succeeds at the state level, but the operator is dense independently.

The invariant is that the cost measures the physical asymmetry of a two-center problem. The currency is a choice; the total is not.

Observation 1 is labeled an observation, not a theorem: it is a pattern across four instances with a plausible cause, not a proof that no fifth currency exists. The underlying obstruction has been characterized in the framework as non-commuting center projections, with $\|[P_A, P_B]\| = 0.50$ at LiH equilibrium and principal angles $7.6^\circ/44.7^\circ/67.3^\circ$.

V. DEMONSTRATION: NAH BINDS WITH GENUINE INTEGRALS

If the wall is the Hamiltonian specification, then supplying genuine end-to-end integrals should bind NaH without any change to the treatment of correlation. It does.

The construction is all-electron 12-electron NaH with $Na\{1s, 2s, 2p \times 3, 3s\}$ and $H\{1s\}$ ($M = 7$ spatial orbitals), non-orthogonal determinants built from fragment-native orbitals, and general- N non-orthogonal Slater–Condon evaluation by Löwdin cofactors. Orbital exponents were optimized on the *isolated* Na atom only ($E(Na) = -161.0845$ Ha against a Clementi minimal-STO reference of ≈ -161.12), with hydrogen at $\zeta = 1$, so the molecular curve is a fragment prediction rather than a fit. Machinery validation: rotational invariance at 1.1×10^{-13} , and the 91-determinant complete space reproduces bitstring FCI with a span identity of exactly zero.

Three determinants give an interior minimum at $+4.8\%$ of the experimental bond length, with the in-basis

TABLE II. NaH, all-electron 12e, $M = 7$. Percentages are of in-basis FCI binding. Experiment: $R_{eq} = 3.566 a_0$, $D_e = 1.961$ eV.

method	$R_{eq} (a_0)$	D_e (eV)	% FCI
covalent (2 det)	3.954	0.686	58.4
+ Na^+H^- (3 det)	3.736	1.071	91.1
all 4 det	3.713	1.080	91.9
FCI (91 det)	3.595	1.175	100
Löwdin covalent (2 det)	—	unbound	−38.0
Löwdin 3 det	3.658	0.136	11.6

FCI landing at +0.8%. The D_e deficit (55% of experiment) is minimal-basis incompleteness—no polarization, no diffuse functions—and reproduces the pattern of the four-electron LiH case at 53%; it is not an integral-fidelity effect.

The Löwdin comparators are the informative rows. Orthogonalizing the orbitals and truncating to the *same* number of determinants retains 11.6% of the binding, and the two-determinant case is destroyed outright. Because the complete space is bit-identically the same either way, the damage is specific to *truncated* spaces. The four-electron LiH case replicates this at 15.0% and destroyed, respectively.

Observation 2 (Orthogonalization damages only truncated spaces). *Orthogonalization is span-preserving and therefore harmless on the complete space, and severely harmful on truncated subspaces built from physically meaningful fragment structures. A qubit register only ever represents a truncated space.*

Observation 2 holds at four and twelve electrons, first and second row, with and without p orbitals. It generalizes to systems not yet run and is cheap to falsify.

VI. SCOPE

What this paper does not claim. It does not claim the framework’s atomic-sector results are affected: single-

center angular structure is untouched, exact, and remains the basis of the atomic qubit results. It does not claim the qubit-resource results of Paper 14 are wrong; it clarifies that they are measured on the builder, whose neglected overlap is the subject here. And it does not claim a route around the obstruction, because the census closed the last one available.

A distinction worth naming explicitly, because it bounds any scaling ambition. *Symmetry sparsity*—exact zeros forced by conservation laws, knowable from labels—is what this framework has, and Theorem 1 shows it is an atomic-sector property that degrades to an abelian residue at two centers and to a finite grading when linearity is lost. *Distance sparsity*—approximate near-zeros from spatial locality, exploited by screening, density fitting, and density-matrix truncation—is what linear-scaling electronic structure runs on, and it *improves* with system size. The framework has the first and none of the second. It is, structurally, a small-system framework, and both of its advantages contract as atom count grows: the abelian grading shrinks with symmetry, while the dense cross-center classes grow with the number of center pairs.

-
- [1] J. Loutey, “Prolate Spheroidal Natural Geometry for Two-Center One-Electron Systems,” Paper 11 in the Geometric Vacuum series.
 - [2] J. Loutey, “Qubit Encodings from Natural Geometry,” Paper 14 in the Geometric Vacuum series.
 - [3] R. S. Mulliken, C. A. Rieke, D. Orloff, and H. Orloff, “Formulas and numerical tables for overlap integrals,” *J. Chem. Phys.* **17**, 1248 (1949).
 - [4] K. Ruedenberg, “A study of two-center integrals useful in calculations on molecular structure,” *J. Chem. Phys.* **19**, 1459 (1951).
 - [5] E. Artacho and L. M. del Bosch, “Nonorthogonal basis sets in quantum mechanics: Representations and second quantization,” *Phys. Rev. A* **43**, 5770 (1991).

TABLE III. Claim $\rightarrow$ backing. Test backfill from the exploratory drivers into **tests/** is a precondition for this paper's release; see Sec. VI.

claim	tier	backing (tests/)
Table I	MEASURED	<code>..p58_census</code>
Eq. (1)	SYMBOLIC	<code>..p58_certificate</code>
Thm. 1	INTERNAL	<code>..p58_m_rule</code>
Pred. 1	UNTESTED	—
Obs. 1	OBSERVATION	CHANGELOG v4.73.x
Table II	PANEL-VER.	<code>..p58_nah</code>
Obs. 2	MEASURED	<code>..p58_nah</code>