

Angular Sparsity Is an Atomic-Sector Property: The Abelian Residue of Multi-Center Bonding*

J. Loutey¹

¹*Independent Researcher, Kent, Washington*
(Dated: August 10, 2026)

The GEOVAC program builds molecular Hamiltonians from atom-centered hydrogenic blocks (screened Z_{eff} per block, *not* shared- p_0 Coulomb–Sturmian, which Papers 8–9 rule out for heteronuclear pairs) whose angular matrix elements are generated from quantum-number labels by Gaunt selection, with no spatial quadrature. This yields structurally sparse qubit Hamiltonians—334 Pauli terms for LiH at $Q = 30$, $O(Q^{2.5})$ scaling (Paper 14)—at the cost of a documented binding failure for second-row diatomics (the W1e wall, Papers 19/20 [4, 5]).

We settle the question of whether that trade was necessary. Using exact rational two-center machinery (Mulliken auxiliary functions A_n/B_n for overlaps; a hydrogenic eigen-trick for cross-center one-body terms, certified by an exact bra/ket identity holding entrywise), we census the *genuine* bare tensor (S, h, g) for a real fragment pair and find: **[MEASURED]** the Gaunt l -selection rule does not survive cross-center evaluation, while the axial m -rule does; cross-center one-body entries measure 0.19–0.80 Ha—bond scale, not perturbative. **[COUNTED]** the symmetry-*permitted* electron-repulsion density inflates $13.8\times$ over the builder at $n_{\text{max}} = 2$ and $15.0\times$ at $n_{\text{max}} = 3$; the $n_{\text{max}} = 2$ figure is independently corroborated numerically (29.8% measured against 29.4% counted), the $n_{\text{max}} = 3$ figure is counted only.

We give the mechanism **[INTERNAL THEOREM]**: m labels the *abelian* symmetry shared by both centers (rotation about the internuclear axis is one global one-parameter group, so m is a common label), while l labels irreducible representations of a rotation group attached to a single center. The intersection of the two centers’ rotation groups is exactly the axial $U(1)$, whose character group is $\mathbb{Z}$ by Pontryagin duality. Abelian selection rules are additive charge conservation and survive; non-abelian Clebsch–Gordan selection requires the full sphere and does not. The surviving *cross-center* angular rule is therefore the one generated by the common axial subgroup, which is readable from the molecular point group in advance.

Two re-readings follow. First, the $17.9\times$ Pauli inflation previously attributed to Löwdin orthogonalization (v4.73.1) was not, in the bulk, orthogonalization’s price: the genuine tensor’s own permitted density inflates by the same order ($\sim 15\times$) before any transformation, so Löwdin was paying a bill that was largely already owed. Second **[OBSERVATION]**, four previously attempted relocations of the non-orthogonality—into the operator, into a dual basis, into a high-symmetry embedding, into device-measured many-body overlaps—each moved the cost without reducing it, because the cost measures the physical asymmetry of a two-center problem. A fifth, measured here, is the sharpest: enabling equivalent-atom-swap tapering *reduces* the qubit saving (BeH_2 $12 \rightarrow 9$, N_2 $22 \rightarrow 19$, F_2 $22 \rightarrow 15$) while inflating the Pauli count 2.3 – $3.6\times$, because the swap rotation merges per-block abelian gradings into joint stabilizers. A permutation identifies blocks, and identifying blocks destroys the per-block labels that generate the sparsity.

We close with a demonstration **[PANEL-VERIFIED]**: all-electron 12-electron NaH, the framework’s hardest documented failure, binds at $R_{\text{eq}} = 3.736 a_0$ (+4.8% vs experiment) with 91.1% of in-basis FCI binding from three fragment-native determinants, once the integrals are genuine end-to-end—with *fewer* correlation degrees of freedom than the runs that failed, which is what makes the inference go through. This localizes W1e to the Hamiltonian specification rather than the correlation treatment on this system. The demonstration deliberately forfeits the sparsity: it makes no sparsity claim.

A by-product bears on the continuous theory the census borrows from: the radial integral at the core of the two-center exchange coupling closes in elementary form (weight one, π -free, no dilogarithm), so the $(AA|BB)$ block is not merely counted but *decided*—all 195 permitted entries proven nonzero by Lindemann–Weierstrass separation.

I. INTRODUCTION: WHAT THE SPARSITY BOUGHT, AND WHAT IT COST

The GEOVAC framework’s computational proposition is structural sparsity. Angular matrix elements are generated from the quantum-number labels (n, l, m) by Gaunt

selection rules and Wigner $3j$ symbols with no spatial integration, so the resulting second-quantized Hamiltonian is sparse by construction rather than by numerical thresholding. The downstream consequence is the qubit-resource result of Paper 14 [2]: composed LiH at 334 Pauli terms on $Q = 30$ qubits, $O(Q^{2.5})$ Pauli scaling against $O(Q^{3.9-4.3})$ for Gaussian baselines, and a $190\times$ reduction against cc-pVDZ.

* Paper 58 in the Geometric Vacuum series

The framework’s documented failure is chemical binding for second-row diatomics. Named W1e, it has resisted an unusually thorough sequence of attacks: six Phillips–Kleinman modifications, basis-level Schmidt orthogonalization, core-correlation estimates, basis enlargement to $n_{\max} = 4$, an explicit-core Hartree–Fock treatment with twelve electrons, and a DMRG cross-check that reproduced FCI bit-identically at finite sector dimension without producing an interior minimum (Papers 19/20 [4, 5]; CHANGELOG v2.19.4–v3.87.0). The wall was localized to the Hamiltonian specification—the projection from the continuous Level-4 solver to second-quantized ($h_1, \text{eri}, e_{\text{core}}$) integrals—rather than to the treatment of correlation.

This paper answers the question that localization leaves open. The framework neglects the overlap between orbitals on different centers, treating a non-orthogonal basis as though it were orthonormal, and that neglect is what preserves the sparsity. Was there an alternative? Could one keep the genuine cross-center physics *and* the label-generated sparsity, by declining to transform the integral tensor at all?

The answer is no, and the reason is group-theoretic rather than technical.

A. A change of method, not a refinement

Relation to the continuous theory. Evaluating integrals over genuine two-center orbital products is, in itself, nothing new: it is standard Coulomb–Sturmian practice, developed over decades by Avery and co-workers on the many-center hyperspherical-harmonic foundation [13, 14], and the framework’s cross-center one-body couplings are precisely the Shibuya–Wulfman integrals [15] of that theory (Paper 19 [4]). GEOVAC adds nothing to that continuous apparatus. What is new is not the integrals but the *discrete* reading of them: which of the continuous theory’s couplings survive as exact, label-forced zeros when the angular content is carried to a qubit Hamiltonian, and which do not. In that sense this paper is the discrete/combinatorial skeleton of a continuous subject—the relationship the angular selection rules already bear to their Gaunt and Wigner origins, now traced into the two-center case.

It is worth stating plainly what this paper does differently, because the step is easy to mistake for an incremental improvement in integral quality. It is not. Every previous GEOVAC molecular builder evaluates its integrals in the *atomic sector*: each orbital product entering an integral sits on a single nucleus, and inter-nuclear coupling enters through a surrogate. Three generations, distinguished by what they do at the bond:

1. **Composed geometry** (Paper 17 [3]). A Phillips–Kleinman pseudopotential replaces explicit two-body core–valence Coulomb with one-body Z_{eff} screening, so *cross-block electron-repulsion integrals*

vanish identically. The bond is not present in the two-body tensor at all.

2. **Balanced coupled composition** (Papers 19/20 [4, 5]). The pseudopotential is eliminated and cross-block repulsion is restored, balanced by a cross-center nuclear-attraction term. But both restored pieces remain atomic-sector objects: the cross-center one-body integral is $\langle \psi_{nlm}^A | -Z_B/|\mathbf{r} - \mathbf{R}_B| | \psi_{n'l'm'}^A \rangle$, with *both* orbitals on center *A* and only the nucleus displaced, and the restored cross-block repulsion is constructed so as to continue obeying Gaunt triangle selection. The orbital products are still one-center; the *l*-selection rule still holds.

3. **This paper.** The integrals are evaluated over orbital products that genuinely span two nuclei, $\langle \chi^A | \cdot | \chi^B \rangle$. That is the step neither prior generation takes.

The distinction is not cosmetic, and the census below measures its price exactly. Because generations 1 and 2 never form a two-center orbital product, the Gaunt *l*-selection rule is exact for them by construction, and the label-generated sparsity follows. Once the product genuinely spans two centers, that rule fails—only the axial *m*-rule survives—and the permitted density inflates by $13.8\times$ at $n_{\max} = 2$. The sparsity was never a discovered property of molecular electronic structure; it was a property of declining to evaluate the bond.

The demonstration in Sec. V is the corresponding gain. NaH is the system the prior generations could not bind under any modification attempted, and it binds here with three fragment-native determinants. We therefore read W1e not as a defect to be repaired within the atomic-sector construction, but as the accumulated cost of that construction—which makes the present route a replacement for it rather than a correction to it. What the replacement forfeits is precisely the qubit-resource proposition Paper 14 rests on; Sec. IV states that trade as a conservation law rather than a defeat.

Scope of this subsection. The three-generation description is a statement about what the builders compute, verifiable from their construction, not a new result. The demonstration below (Sec. V) is a diatomic. A polyatomic extension has since been measured (BeH_2 and H_2O , basis ladder to $M = 19$): at matched minimal basis and matched correlation treatment the present route improves the geometry only modestly over generation 1 (+10.0% vs 11.7% for BeH_2 , +14.9% vs 19.4% for H_2O), but the error is basis-limited rather than method-limited and falls to +1.4% and +2.5% respectively once the basis is enlarged—and the H_2O bond angle, an observable generation 1 cannot represent at all, comes out at 106.3° against an experimental 104.5° . Those figures are RHF over Gaussian-fitted Slater shapes in the same engine used for Sec. V; they are reported here as scope, not as a result of this paper.

II. THE EXACT-ARITHMETIC CENSUS

A. Machinery and certificate

Two-center overlaps at arbitrary (n, l, m, Z) are computed as exact rationals via the Mulliken auxiliary functions A_n/B_n [7, 8]—one of several classical routes to two-center Slater integrals [19, 20]—evaluated in prolate spheroidal coordinates where the two-center problem separates [1]. Because the result is rational, vanishing is *decidable* rather than inferred from a numerical tolerance—a property no quadrature scheme can offer, and the reason this census can make a claim about exact zeros at all.

Cross-center one-body elements are obtained by extending the same route with $1/r_A$ and $1/r_B$ kernels, each of which cancels against the volume element $(\xi^2 - \eta^2)$, followed by the hydrogenic eigen-trick. The assembly carries its own certificate: the bra/ket identity

$$E_a S - Z_B I_B = E_b S - Z_A I_A \quad (1)$$

holds as an *exact rational equality* on a six-case panel and on every entry of the census configuration—and coefficient-wise on both exponential legs, which is strictly stronger than agreement at a single R . Its diagnostic power is itself tested rather than assumed: two independent corruptions of the inputs (a wrong bra eigenvalue, and I_A interchanged with I_B) each break the identity, so it is demonstrably not an algebraic tautology. Four independent quadrature cross-checks agree at $\sim 2 \times 10^{-16}$.

B. Result

The census configuration is $Z_A = 3$ at the origin, $Z_B = 1$ at $R = 3\hat{z}$, hydrogenic orbitals to $n_{\max} = 2$ on each center, $M = 10$ spatial orbitals. Table I compares the genuine bare tensor against the production builder’s.

Provenance, stated precisely because the three rows differ, and because the rows are internally mixed. The cross-center blocks of S and h (22 entries each) are evaluated in exact rational arithmetic by the machinery of Sec. II A, so their *non*-vanishing is decided rather than inferred. The same-center S block is fixed analytically by orthonormality of same- Z hydrogenics rather than recomputed; the same-center h block is evaluated with the production

TABLE I. Genuine versus builder sparsity, $Z_A = 3/Z_B = 1$ at $R = 3a_0$, $n_{\max} = 2$ per center ($M = 10$). “Genuine zeros are” records which selection rule accounts for the surviving zeros.

object	genuine	builder	dense	genuine zeros are
S	32	10	100	m -rule + same-center
h	44	22	100	m -rule only
g	2,944	214	10,000	m -rule + same-center Gaunt

float routine, so its zeros there are threshold decisions. And the vanishing of every non-matching- m entry is not evaluated at all—it is forced by Theorem 1(i). In this census the rational machinery decides non-vanishing; the zeros come from symmetry. The g row is *counted*: it is a symmetry-rule census in the complex- m basis, because no two-center electron-repulsion engine exists in the framework—the m -rule and per-side Gaunt feasibility are applied combinatorially, and no two-center electron-repulsion integral was numerically evaluated. The g figures are therefore claims about how many entries the selection rules *permit*, not verifications that each permitted entry is nonzero. Supplying that engine is the natural follow-on to this paper.

Status of that follow-on. It is now partially supplied, and the per-class outcome bears directly on what the g row could become. Three of the four classes are closed in exact form and validated against independent engines: the one-center classes, $(AA|BB)$, and the hybrid pair $(AA|AB)/(AB|BB)$. For the exchange class $(AB|AB)$ —Ruedenberg’s own subject [9]—the assembly is validated for general (l, m) and its ordered radial double integral, the last non-closed step, now closes as well, for general (τ, σ, H, j) and unequal exponents (agreement with nested quadrature at 10^{-13}). All four classes therefore have closed-form machinery.

class	closed form	transcendental content
one-center, $(AA BB)$	yes	none—elementary
$(AA AB)$, $(AB BB)$	yes	$\{E_1, \ln\}$
$(AB AB)$	yes	$\{E_1, \ln, \gamma\}$

The transcendental content grows across the classes, but—and this is the structurally important part—it does *not* escalate in weight. The ordered radial integral is an iterated integral over a simplex, the shape that defines a period, and iterated integrals of weight-one objects generically land at weight two (dilogarithms, $\zeta(2)$). This one does not: assembled in closed form it agrees with quadrature to 6.3×10^{-16} and contains only exp, E_1 , $\ln$ and Euler’s γ . Nor does $\sigma \neq 0$ break it. The derivatives $d^\sigma Q_\tau / d\xi^\sigma$ carry poles of order up to σ at $\xi = \pm 1$, but they are absorbed exactly by the $(\xi^2 - 1)^H$ prefactor, because per electron the net exponent there is $(|m_a| + |m_b| - |m_a - m_b|)/2 \geq 0$ by the triangle inequality—with equality, i.e. exact absorption, whenever m_a and m_b have opposite signs. Verified to 2.5×10^{-14} with no weight-two content at any σ tested. So the seed sets above form a weight-one filtration throughout.

This does not upgrade the g row, for two independent reasons. First, no re-census has been run with the new machinery, so the counted figures stand exactly as counted. Second, and structurally, what the row needs is zero-*decidability*, and that is available only where a closed form exists.

For $(AA|BB)$ that argument now goes through, and by exactly the route the S and h rows already use. Every π

cancels from the closed form—the harmonic normalisations against the $4\pi/(2L+1)$ of the multipole potential—leaving

$$(ab|cd) = A_0(R) + \sum_j A_j(R) e^{-\lambda_j R},$$

with A_j rational functions and λ_j rational. Checked on $(1s\,1s|1s\,1s)$, $(2p_0\,2p_0|1s\,1s)$ and $(2p_{+1}\,2p_{+1}|2p_0\,2p_0)$ across Z combinations: the π power is common (indeed zero) in every term. For algebraic R the exponents $-\lambda_j R$ are distinct algebraic numbers, so by Lindemann–Weierstrass the family $\{1, e^{-\lambda_j R}\}$ is linearly independent over the algebraic numbers and the expression vanishes iff every A_j does. That is the same separation the S and h rows are decided by, applied to a class that now has the closed form to support it.

The hybrid and exchange classes do not follow: they carry $\{E_1, \ln\}$ and $\{E_1, \ln, \gamma\}$ respectively, and would need an independence argument over those seeds, which we have not attempted.

Consequence for this paper’s own sector. The classes carrying residual l -selection— $(AA|AA)$, $(BB|BB)$, $(AA|BB)$ —are precisely the ones that are now zero-decidable: the one-centre classes trivially—they are exact rationals—and $(AA|BB)$ by the argument above. So the sector that Theorem 1 makes a positive prediction about is no longer merely countable but decidable.

And the decision has now been run on $(AA|BB)$. On the census configuration, that block holds 625 entries of which the counting rules permit 195. Deciding each one by the separation above—grouping the closed form by exponential rate and asking whether every rational coefficient vanishes—gives

outcome	count
decided nonzero	195
zero, vanishing Gaunt coefficient	0
zero, accidental (Gaunt nonzero, radial sum cancels)	0

All 195 permitted entries are genuinely nonzero.

There are no accidental zeros and no symmetry zeros that the counting missed, so on this block the counted density is not an overcount but is *exactly* the true density. This is a decision rather than a measurement: no threshold is involved anywhere, and it upgrades the block from COUNTED to DECIDED.

It also settles, for this block, the reading the Gaussian corroboration could only suggest. That check found 29.8% against 29.4% counted and attributed the gap to the real-Cartesian versus complex- m basis convention rather than to over-permissive counting; the decided census confirms there is no over-permissiveness here to find.

The g row’s overall tier is nonetheless unchanged, because $(AA|BB)$ and the one-centre classes together are only the 20.5% that carries residual l -selection. The three genuinely cross classes—80% of the tensor—remain

counted, and will until an independence argument over $\{E_1, \ln, \gamma\}$ is available.

We state the outlook carefully, because the natural pessimistic reading is wrong. The obstruction is *not* that the seed set escalates: it is a weight-one filtration at every class, so the exchange case is a harder independence question than the hybrid case but not a categorically different one. Nor is it that a closed form is missing: one now exists for every class. What is missing is the independence arguments themselves, plus a re-census run through the new machinery. Both are ordinary work rather than obstructions in principle—which is a materially better position than the counted row has occupied until now, and materially worse than a decided row.

What the build does settle is narrower, and worth stating because it is exactly the sector this paper’s mechanism concerns. The classes in which the residual l -selection survives— $(AA|AA)$, $(BB|BB)$ and $(AA|BB)$, together 20.5% of the genuine tensor—are now exactly computable in closed form and carry no transcendental seed whatever. An l -selection census restricted to those classes is therefore no longer gated on the missing engine. That is the sharpest thing the build supplies for this paper: the sector where Theorem 1 predicts residual l -selection to survive is now exactly computable, and computable with no transcendental seed at all, so a census over it would test the theorem’s positive content directly rather than by counting.

The two halves compose. The census machinery of Sec. II supplies exact two-centre S and h —the Mulliken/Ruedenberg auxiliary route and the hydrogenic eigen-trick—and named the missing two-centre electron-repulsion engine as its one gap. With that engine supplied, the three ingredients now compose into a complete calculation carrying no Gaussian fit anywhere. Verified on H_2 at $R = 1.4$ bohr in a $1s$ -per-centre basis: total FCI energy -1.106556606091 Ha. Evaluating the identical basis through Gaussian-fitted integrals converges monotonically onto that value as the fit deepens (2.7×10^{-3} , 3.0×10^{-4} , 4.5×10^{-5} , 8.0×10^{-6} , 1.9×10^{-6} Ha at 4, 6, 8, 10, 12 primitives), confirming the native value is the exact one. The residual sits in h , not in g : the electron-repulsion tensors agree to 1.9×10^{-8} and the overlaps to 1.5×10^{-8} at ten primitives, while the one-electron terms lag because kinetic energy probes the nuclear cusp, where a sum of Gaussians converges slowly by construction.

Subject to that distinction, the electron-repulsion tensor is 29.4% permitted-dense genuinely against 2.1% for the builder.

Numerical corroboration of the g row. The counting can be checked without the missing Neumann engine [9], by evaluating the same tensor with a McMurchie–Davidson Gaussian engine [11] over fitted Slater shapes. On the census configuration this gives 2,976 of 10,000 entries above 10^{-10} , i.e. 29.8% dense against the 29.4% counted—agreement at the level the basis convention permits, since the numerical evaluation is in the real

Cartesian basis and the count is in the complex- m basis, which differ by the $\pm m$ mixing. So the permitted entries are *generically nonzero*: the density is a real property of the tensor and not an artifact of over-permissive counting. The l -selection failure appears at bond scale in individual entries—for instance $(1s_A 2p_{z,B} | 2s_A 2s_A) = -0.218$, where the same-center analog vanishes identically.

This does not upgrade the row to MEASURED: Gaussian-fitted floats make these threshold decisions rather than decided zeros, so the exact statement still awaits the engine. What it removes is the possibility that the counted density was an overcount.

The m -rule, verified without reference to m . Theorem 1(i) has a basis-free consequence that can be tested directly on the tensor: since $1/r_{12}$ is rotationally invariant and both nuclei lie on the axis, a rotation about the internuclear axis maps the basis into itself and must leave the tensor fixed. It does, to 2.2×10^{-16} . Displacing the nuclei off that axis and repeating the identical check gives 1.3×10^{-1} —fifteen orders of magnitude larger. The symmetry is axial, and only axial. At $n_{\max} = 3$ ($M = 28$) the comparison is 114,280 *permitted* entries (18.6%) against 7,600 (1.2%): the inflation factor is $13.8\times$ at $n_{\max} = 2$ and $15.0\times$ at $n_{\max} = 3$, i.e. *increasing* with basis size. Both are counted; only the $n_{\max} = 2$ figure is numerically corroborated below.

Two further measurements matter for interpretation. First, per-side Gaunt feasibility removes 5.6% of entries at $n_{\max} = 2$ and 6.3% at $n_{\max} = 3$, beyond what the m -rule already removes, and the three cross classes $(AA|AB) + (AB|AB) + (AB|BB)$ account for 80% of the genuine tensor. The residual l -selection is therefore confined to the classes carrying a same-center distribution on each side— $(AA|AA)$, $(BB|BB)$ and $(AA|BB)$ —together 20.5% of the genuine tensor; the three genuinely cross classes carry none. Second, cross-center magnitudes are large: overlaps 0.08–0.52, one-body entries 0.19–0.80 Ha. For scale, the entire NaH bond is under 2 eV, so individual entries in the dense part of the tensor exceed the binding energy being computed. They are not droppable as small corrections—an independent quantification of the W1d/W1e diagnosis.

III. MECHANISM: THE ABELIAN RESIDUE

The census result is not an accident of the chosen configuration. It is forced, and the forcing is elementary once stated in the right language.

Theorem 1 (Angular selection survives molecularization only in the abelian sector). *Let two nuclei A, B lie on the z -axis, with atom-centered basis functions labeled (n, l, m) about their own center, and let $\hat{O}$ be invariant under the molecular symmetry group. Then:*

- (i) *The m -selection rule holds exactly across centers:*
 $\langle ab | \hat{O} | cd \rangle = 0$ unless $m_a + m_b = m_c + m_d$.

- (ii) *No l -selection rule holds between basis functions on different centers. More strongly, the cross-center support is full in l at fixed m : the re-expansion does not terminate, so every admissible l couples.*

- (iii) *Of the per-center $SO(3)$ labels, only those of the common axial subgroup $U(1) \cong SO(2)$ survive cross-center; by Pontryagin duality its characters are the integers, so the surviving cross-center rule is additive m -conservation.*

Part (iii) is deliberately narrow, and two limits are worth stating because the wider reading is false. It does not say the axial $U(1)$ generates all surviving sparsity: same-center Gaunt selection survives from the full per-center $SO(3)$ (visible in Table I's g and S rows), and the largest group fixing both nuclei is not $U(1)$ but $C_{\infty v}$, which is *non-abelian*—its σ_v reflection is exactly the Hopf $m_l \rightarrow -m_l \mathbb{Z}_2$ used as the baseline of Table II. So a non-abelian common group does yield usable label structure. What (iii) claims is only that the *continuous* per-center angular labels reduce to the axial one across centers, which is what the census measures.

Proof sketch. For (i): rotation about the internuclear axis is a *single* one-parameter group acting on all of $\mathbb{R}^3$. Its generator L_z is independent of which point on the axis is taken as origin, so m is a global label shared by both centers rather than a per-center one. Invariance of $\hat{O}$ under this group gives additive conservation of m .

For (ii): l labels irreducible representations of the rotation group about a *specific* center. Rotations about A and rotations about B are distinct subgroups of the Euclidean group whose intersection is precisely the axial subgroup. A function that is an l -eigenfunction about B , re-expanded about A , has support on all l : the two-center re-expansion does not terminate. Hence no l -selection between cross-center pairs. (The same-center blocks retain Gaunt selection, consistent with the $< 6\%$ residual measured in Sec. II.)

For (iii): the intersection in (ii) is $SO(2) \cong U(1)$, abelian. Abelian compact groups have one-dimensional irreducibles labeled by characters, and by Pontryagin duality the character group of $U(1)$ is $\mathbb{Z}$ —which is the integer label m . Selection rules from abelian symmetry are additive conservation laws on integer charges, and are preserved by any operation respecting the axis. Non-abelian selection rules are Clebsch–Gordan couplings requiring the full group; reducing the group destroys them. $\square$

The theorem is independently corroborated by the framework's own production code, along a path that knows nothing of this analysis: the cross-center nuclear-attraction routine in `shibuya.wulfman` returns zero by an explicit early exit when $m_1 \neq m_2$, and imposes no l -selection on the element: its per-multipole Gaunt conditions (triangle inequality and $l_1 + L + l_2$ parity) are

satisfiable by some L for every (l_1, l_2) , so no l pair is excluded. The m -rule is hard-coded because it holds; the l -rule is absent because it does not.

Theorem 1 converts a measurement into a prediction, because the largest common abelian symmetry is readable from the molecular point group before any integral is computed.

Corollary 1 (Sparsity budget from the point group). *For a molecule with point group G , label-generated angular sparsity is bounded by the grading available from the abelian structure of G . Linear molecules ($C_{\infty v}$, $D_{\infty h}$) carry a continuous $U(1)$ and therefore a $\mathbb{Z}$ grading, worth a multiplicative factor. Bent molecules carry only a finite point group and therefore a finite grading.*

H_2O has point group C_{2v} , which is order four and abelian, isomorphic to $\mathbb{Z}_2 \times \mathbb{Z}_2$; all its irreducibles are one-dimensional, so Corollary 1 predicts a two-bit spatial grading—two qubits, not a multiplicative factor. We attempted to test this and could not, for a reason that is itself a structural finding.

Observation 1 (The composed builder carries no angular geometry). *`molecular_spec.py` contains no bond angle: `h2o_spec` is `hydride_spec(8)`, parameterized by nuclear charge and a single O-H distance distributed across five blocks (core, two bond pairs, two lone pairs), with `MolecularSpec.nuclei = None`. The composed H_2O Hamiltonian therefore does not represent C_{2v} at all, and the framework’s reported H_2O geometry error is a purely radial coordinate with no angular degree of freedom. Corollary 1 is consequently not well posed on this builder in the bent case; supplying hand-built bent nuclei would test a symmetry the Hamiltonian does not encode.*

Prediction 1 (Conditional, for any angular build). *In a build that does carry angular geometry, a molecule whose point group is abelian of order 2^k admits a k -bit spatial grading and no more. This remains falsifiable, but not on the present builder.*

What is measurable on the present builder is the linear case, where geometry is angle-free and therefore unambiguous—and it returns a result we did not anticipate, reported as Observation 3 below.

The measured factor of 3–5 for the m -rule at the bases of Sec. II is the $\mathbb{Z}$ -grading case of Corollary 1. It is exact and structural, and it is a constant factor—not the $O(Q^{2.5})$ scaling story the framework’s qubit results rest on.

IV. TWO RE-READINGS

A. Löwdin was paying an honest bill

The corpus previously recorded a $17.9\times$ Pauli inflation from retrofitting Löwdin orthogonalization [10] onto the balanced LiH integrals (CHANGELOG v4.73.1), and a

$14\times$ inflation in the earlier heterogeneous-nested experiment. Both were read as the price of orthogonalization, and orthogonalization was accordingly treated as the sparsity-destroying move.

The census shows this reading was wrong in an instructive way. The genuine tensor’s own permitted-density inflation over the builder is $\sim 15\times$ —the same order as the $17.9\times$ Löwdin figure, though the two are not commensurable: $17.9\times$ is a Pauli-term count on Löwdin-retrofitted balanced LiH integrals, $15.0\times$ a counted permitted-entry ratio on a $Z_A = 3/Z_B = 1$ hydrogenic census, and neither carries an error bar. The inference does not need them to agree numerically: the *bulk* of that inflation cannot have been orthogonalization’s price, because the untransformed tensor never possessed the sparsity Löwdin was said to destroy. “Löwdin destroys sparsity” and “the molecular tensor is intrinsically l -dense cross-center” are the same fact observed twice, from two sides.

This matters because it removes an entire class of proposed remedies. Any scheme whose premise is “avoid the transformation and keep the sparsity” is answering a question that was mis-posed.

B. Cost conservation

Observation 2 (Relocating the metric does not reduce it). *Five independent attempts to move the non-orthogonality out of the way each succeeded in moving it and failed to reduce it:*

1. into the **operator** (Löwdin $S^{-1/2}$): operator densifies $17.9\times$;
2. into a **dual basis** (biorthogonal encoding): densifies comparably. The metric provably does not disappear—in the contravariant representation it reappears as an S^{-1} contraction, and the dual creation/annihilation operators cease to be Hermitian conjugates [12]—but whether that S^{-1} is dense is a property of the two-center Slater basis, measured here rather than imported;
3. into a **high-symmetry embedding** (hyperspherical Level $4N$): angular basis requirement explodes—LiH at 63.5% R_{eq} error with $l_{\text{max}} = 2$ genuinely insufficient;
4. into **device-measured many-body overlaps** (NOCI): succeeds at the state level, but the operator is dense independently;
5. into a **permutation symmetry** (equivalent-atom swap): net loss, measured—see Observation 3.

The invariant is that the cost measures the physical asymmetry of a two-center problem. The currency is a choice; the total is not.

The fifth entry was measured for this paper and is the sharpest instance, because the trade is visible term by term.

Observation 3 (Permutation symmetry costs qubits on this builder). *Enabling equivalent-atom-swap tapering reduces the qubit saving and inflates the Pauli count, in every swap-eligible configuration tested (Table II). The mechanism is visible in the retained stabilizer labels: the swap rotation merges the per-block Hopf and ℓ -parity stabilizers into joint **SWAP-JOINT** stabilizers, so one permutation stabilizer is bought at the price of several per-block abelian gradings. [MEASURED]*

This is consistent with Theorem 1 rather than incidental to it. A permutation is not an abelian grading of the existing block structure; it is an operation that *identifies* blocks, and identifying blocks destroys the per-block labels that generated the sparsity. The abelian residue is not merely what survives—it is what can be exploited without collapsing the structure that carries it.

Documentation discrepancy. The docstring of `extended_tapered_from_spec` states that atom-swap “saves additional qubits on polyatomics ($\text{BeH}_2 + 1$, $\text{H}_2\text{O} + 1$, $\text{NH}_3 + 1$, $\text{CH}_4 + 2$).” The measured BeH_2 value is -1 against the Hopf-only baseline and -3 against Hopf + ℓ ; the sign is wrong under both. We report the measurement and do not attempt to adjudicate whether the documentation was always incorrect or the behavior regressed.

Observation 2 is labeled an observation, not a theorem: it is a pattern across five instances with a plausible cause, not a proof that no fifth currency exists. The underlying obstruction has been characterized in the framework as non-commuting center projections, with $\| [P_A, P_B] \| = 0.50$ at LiH equilibrium and principal angles $7.6^\circ/44.7^\circ/67.3^\circ$.

V. DEMONSTRATION: NAH BINDS WITH GENUINE INTEGRALS

If the wall is the Hamiltonian specification, then supplying genuine end-to-end integrals should bind NaH *without needing more* correlation than the failing runs used. It does—with considerably less.

Two scope conditions govern what follows, and the argument is weaker without them. First, this construction **makes no sparsity claim**: the radial functions

TABLE II. Equivalent-atom-swap tapering against two baselines. ΔQ is qubits removed (larger is better); Pauli counts are for the tapered operator. LiH has no equivalent atoms and is the null control.

system	baseline	ΔQ	ΔQ +swap	Pauli ratio
BeH_2	Hopf only	7	6	$3.50\times$
BeH_2	Hopf + ℓ	12	9	$3.60\times$
N_2	Hopf + ℓ	22	19	$2.30\times$
F_2	Hopf only	12	9	$3.50\times$
F_2	Hopf + ℓ	22	15	$3.60\times$
LiH	Hopf + ℓ	8	8	$1.00\times$

are six-Gaussian fits of Slater shapes evaluated with a McMurchie–Davidson engine [11] and a dense non-orthogonal Slater–Condon solver, so it forfeits precisely the label-generated sparsity this paper is about—it is item 4 of Observation 2, not an escape from it. Second, the correlation treatment is not held fixed: the framework’s failing runs were FCI and DMRG in an orthogonal second-quantized space, whereas this is three non-orthogonal determinants. That asymmetry is what makes the inference go through—binding appears with *less* correlation, so correlation cannot have been what was missing.

The construction is all-electron 12-electron NaH with $\text{Na}\{1s, 2s, 2p\times 3, 3s\}$ and $\text{H}\{1s\}$ ($M = 7$ spatial orbitals), non-orthogonal determinants built from fragment-native orbitals [17, 18], and general- N non-orthogonal Slater–Condon evaluation by Löwdin cofactors [16]. Orbital exponents were optimized on the *isolated* Na atom only ($E(\text{Na}) = -161.0845$ Ha against a Clementi minimal-STO reference of ≈ -161.12), with hydrogen at $\zeta = 1$, so the molecular curve is a fragment prediction rather than a fit. Machinery validation at $R = 3.5a_0$: rotational invariance at 1.1×10^{-13} , and the 91-determinant non-orthogonal complete space reproduces the Löwdin bitstring FCI to 1.8×10^{-15} Ha. (The value serialized by the exploratory driver was corrupted by variable shadowing—the record was written after a geometry sweep that rebinds the nuclear-repulsion term—and has been corrected; the offset was exactly $11/3.5 - 11/10 = 143/70$.)

Three determinants give an interior minimum at $+4.8\%$ of the experimental bond length, with the in-basis FCI landing at $+0.8\%$. The D_e shortfall—55% of experiment retained, so a 45% deficit—is minimal-basis incompleteness (no polarization, no diffuse functions, and no basis-set-superposition correction, the atom references being computed in atom-only bases) and reproduces the pattern of the four-electron LiH case at 53%; it is not an integral-fidelity effect.

As with LiH, the reading is sharper at *fixed* geometry, where no dissociation scan enters and the numbers

TABLE III. NaH, all-electron 12e, $M = 7$. Percentages are of in-basis FCI binding. Experiment: $R_{\text{eq}} = 3.5667a_0$ ($r_e = 1.8874$ Å) [6]. The comparison value $D_e = 1.961$ eV is inherited from the framework’s earlier compilation and its primary source is *not yet verified*; it is not tabulated on the NIST pages carrying r_e , so every D_e percentage below should be read as provisional pending that check.

method	$R_{\text{eq}} (a_0)$	D_e (eV)	% FCI
covalent (2 det)	3.954	0.686	58.4
+ Na^+H^- (3 det)	3.736	1.071	91.1
all 4 det	3.713	1.080	91.9
FCI (91 det)	3.595	1.175	100
Löwdin covalent (2 det)	—	unbound	-38.0
Löwdin 3 det	3.658	0.136	11.6

are pinned by test. At $R = 3.5 a_0$, against an isolated-atom dissociation reference of -161.584337 Ha and an in-basis FCI binding of 1.170 eV: two covalent determinants recover 50.9% of that binding, three recover 89.3%, and all four 90.6%—while the *same* three on Löwdin-orthogonalized orbitals recover 10.8%, and the Löwdin two-determinant case is not merely unbound but actively repulsive relative to the separated atoms (-153%). The well-minimum figures in Table III and these fixed-geometry figures tell the same story by two routes.

The Löwdin comparators are the informative rows. Orthogonalizing the orbitals and truncating to the *same* number of determinants retains 11.6% of the binding, and the two-determinant case is destroyed outright. Because the complete space is bit-identically the same either way, the damage is specific to *truncated* spaces. The four-electron LiH case replicates this at 15.0% and destroyed, respectively.

Observation 4 (Orthogonalization damages only truncated spaces). *Orthogonalization is span-preserving and therefore harmless on the complete space, and severely harmful on truncated subspaces built from physically meaningful fragment structures. A qubit register only ever represents a truncated space.*

Observation 4 holds at four and twelve electrons, first and second row, with and without p orbitals—the four-electron leg pinned by test, the twelve-electron leg on the exploratory driver. We expect it to generalize, and it is cheap to falsify.

The mechanism is also visible at *fixed* geometry, without reference to a dissociation limit, which makes it the cleanest form to check. On LiH at $R = 3.25 a_0$ in a three-orbital fitted-Slater basis, against a complete-space FCI of -8.887049 Ha: three fragment-native determinants land $+1.32$ mHa above it, while the *same* three determinants built on Löwdin-orthogonalized orbitals land $+40.52$ mHa above—a $30.6\times$ larger error at identical determinant count. At two determinants the errors are $+19.4$ mHa and $+118.7$ mHa ($6.1\times$). The complete spaces agree to machine precision, so the damage is attributable to truncation and not to the orthogonalization corrupting the integrals.

VI. SCOPE

What this paper does not claim. It does not claim the framework’s atomic-sector results are affected: single-

center angular structure is untouched, exact, and remains the basis of the atomic qubit results. It does not claim the qubit-resource results of Paper 14 are wrong; it clarifies that they are measured on the builder, whose neglected overlap is the subject here. And it does not claim a route around the obstruction: the census closes the one route the roadmap had left—untransformed bare-tensor use—while Observation 2 records only that the five relocations attempted so far each moved the cost without reducing it.

A distinction worth naming explicitly, because it bounds any scaling ambition. *Symmetry sparsity*—exact zeros forced by conservation laws, knowable from labels—is what this framework has, and Theorem 1 shows it is an atomic-sector property that degrades to an abelian residue at two centers and to a finite grading when linearity is lost. *Distance sparsity*—approximate near-zeros from spatial locality, exploited by screening, density fitting, and density-matrix truncation—is what linear-scaling electronic structure runs on, and it *improves* with system size. The framework has the first and none of the second. It is, structurally, a small-system framework, and both of its advantages contract as atom count grows: the abelian grading shrinks with symmetry, while the dense cross-center classes grow with the number of center pairs.

TABLE IV. Claim $\rightarrow$ backing. LIVE = passing test in the `tests/test_paper58.*` suite (four files; per-claim mapping in `docs/claim_test_matrix.md`). OWED = result stands on an exploratory driver and must be migrated self-contained into `tests/` before this paper is a finished artifact. CORROB. = numerically corroborated but not exactly decided; see Sec. VI.

claim	tier	backing
Table I S, h	MEASURED	LIVE
Table I g	COUNTED	CORROB.
2-center ERI engine, per class	MEASURED	LIVE
axial invariance	MEASURED	LIVE
Eq. (1)	SYMBOLIC	LIVE
Thm. 1(i)	INT.	LIVE
Thm. 1(ii)	INT.	LIVE
2-center machinery	SYMB.	LIVE
Obs. 1	MEASURED	LIVE
Pred. 1	UNTESTABLE	(needs angular build)
Obs. 2	OBSERVATION	CHANGELOG v4.73.x
Obs. 3	MEASURED	LIVE
Tab. III, fixed R	MEAS.	LIVE
Tab. III, minima	PANEL	OWED
NaH atom-only ζ	MEAS.	LIVE
Obs. 4	MEASURED	LIVE
fit quality, span id.	MEASURED	LIVE

- [1] J. Loutey, “The Molecular Fock Projection: Diatomic Molecules as Prolate Spheroidal Lattices,” GeoVac Paper 11 (2026).
- [2] J. Loutey, “Structurally Sparse Qubit Hamiltonians from

- Spectral Graph Theory,” GeoVac Paper 14 (2026).
- [3] J. Loutey, “Composed Natural Geometries for Core-Valence Diatomics: LiH and BeH⁺ from Discrete Graph Topology,” GeoVac Paper 17 (2026).

- [4] J. Loutey, “Toward PK-Free Molecular Hamiltonians: Two-Center Sturmian Integrals and Coupled Composition,” GeoVac Paper 19 (2026).
- [5] J. Loutey, “Structurally Sparse Molecular Hamiltonians for Near-Term Quantum Simulation,” GeoVac Paper 20 (2026).
- [6] K. P. Huber and G. Herzberg, *Constants of Diatomic Molecules* (Van Nostrand Reinhold, New York, 1979); $^{23}\text{NaH } X^1\Sigma^+ r_e = 1.8874 \text{ \AA} = 3.5667 a_0$, as tabulated in the NIST Chemistry WebBook (accessed 2026-08-11).
- [7] R. S. Mulliken, C. A. Rieke, D. Orloff, and H. Orloff, “Formulas and numerical tables for overlap integrals,” *J. Chem. Phys.* **17**, 1248–1267 (1949).
- [8] C. C. J. Roothaan, “A study of two-center integrals useful in calculations on molecular structure. I,” *J. Chem. Phys.* **19**, 1445–1458 (1951).
- [9] K. Ruedenberg, “A study of two-center integrals useful in calculations on molecular structure. II. The two-center exchange integrals,” *J. Chem. Phys.* **19**, 1459–1477 (1951).
- [10] P.-O. Löwdin, “On the non-orthogonality problem connected with the use of atomic wave functions in the theory of molecules and crystals,” *J. Chem. Phys.* **18**, 365–375 (1950).
- [11] L. E. McMurchie and E. R. Davidson, “One- and two-electron integrals over Cartesian Gaussian functions,” *J. Comput. Phys.* **26**, 218–231 (1978).
- [12] E. Artacho and L. Miláns del Bosch, “Nonorthogonal basis sets in quantum mechanics: Representations and second quantization,” *Phys. Rev. A* **43**, 5770–5777 (1991).
- [13] J. Avery, *Hyperspherical Harmonics: Applications in Quantum Theory* (Kluwer Academic, Dordrecht, 1989).
- [14] J. E. Avery and J. S. Avery, *Generalized Sturmians and Atomic Spectra* (World Scientific, Singapore, 2006).
- [15] T. Shibuya and C. E. Wulfman, “Molecular orbitals in momentum space,” *Proc. R. Soc. Lond. A* **286**, 376 (1965).
- [16] P.-O. Löwdin, “Quantum theory of many-particle systems. I. Physical interpretations by means of density matrices, natural spin-orbitals, and convergence problems in the method of configurational interaction,” *Phys. Rev.* **97**, 1474–1489 (1955).
- [17] A. J. W. Thom and M. Head-Gordon, “Hartree–Fock solutions as a quasidiabatic basis for nonorthogonal configuration interaction,” *J. Chem. Phys.* **131**, 124113 (2009).
- [18] E. J. Sundstrom and M. Head-Gordon, “Non-orthogonal configuration interaction for the calculation of multielectron excited states,” *J. Chem. Phys.* **140**, 114103 (2014).
- [19] M. P. Barnett and C. A. Coulson, “The evaluation of integrals occurring in the theory of molecular structure. Parts I & II,” *Philos. Trans. R. Soc. Lond. A* **243**, 221–249 (1951).
- [20] E. Filter and E. O. Steinborn, “Extremely compact formulas for molecular two-center one-electron integrals and Coulomb integrals over Slater-type atomic orbitals,” *Phys. Rev. A* **18**, 1 (1978).