

Wilson $U(1)$ Lattice Gauge on the Dirac-Rule-B Graph: Structural Compatibility with 3D Compact $U(1)$ on a Non-Cubic Substrate, Including the Polyakov Monopole-Density Signature

GeoVac Project

May 2026

Abstract

The Hopf-graph Wilson $U(1)$ construction of Paper 25 has a known structural limitation: the scalar Fock graph factorises as a disjoint union of 2D rectangular grids $P_{n_{\max}-\ell} \times P_{2\ell+1}$, one per orbital angular momentum ℓ , so its Wilson lattice gauge content reduces per- ℓ to 2D compact $U(1)$ with no genuine 3D phase structure. The Dirac-Rule-B graph of Paper 29 §RH-C supplies a structurally different substrate: dipole edges $\Delta\ell = \pm 1$ connect adjacent angular-momentum shells via E1 parity flips, producing a 1-component graph whose every edge is cross- ℓ and whose every plaquette spans at least two ℓ -shells. This paper specifies $U(1)$ Wilson lattice gauge theory on Rule B, verifies the Hodge identity at machine precision, and tests structural compatibility with 3D compact $U(1)$ via seven independent diagnostics, four leading-order, one non-perturbative monopole-density signature, one full-MC Wilson-loop diagnostic, and one finite-temperature Polyakov-loop diagnostic. (i) Spectral dimension: heat-kernel d_s climbs monotonically from 1.86 at $n_{\max} = 3$ to 2.54 at $n_{\max} = 5$ while the scalar Fock graph plateaus near 1.8. (ii) Migdal–Kadanoff bond-moving RG: β_{eff} flows monotonically to $\beta = 0$ at every tested cutoff with no non-trivial fixed point, consistent with the Polyakov–Banks–Myerson–Kogut permanent confinement mechanism. (iii) Gaussian Wilson loop on the co-exact Hodge sector $K = d_1^\top d_1$: scaling exponent $\alpha \in [0.89, 1.18]$ across $n_{\max} \in \{2, 3, 4, 5\}$ – the free-photon perimeter law that any 3D abelian theory gives at the linearised level. (iv) Leading-order strong-coupling character expansion: $\langle W(C) \rangle = (I_1(\beta)/I_0(\beta))^{A(C)}$ gives an area law with string tension $\sigma(\beta) > 0$ at every tested β , monotonically vanishing toward zero at large β with the correct asymptotics $\sigma \rightarrow \ln(2/\beta)$ and $\sigma \rightarrow 1/(2\beta)$; the perimeter-vs-area discriminator is satisfied at $n_{\max} = 2, 3, 4$. (v) *Non-perturbative Polyakov monopole-density signature* (Sprint XCWG-F, May 2026): a DeGrand–Toussaint construction on Rule B’s size-3 triangle-prism elementary 2-cycles, combined with a Metropolis–Hastings Monte Carlo of compact $U(1)$ Wilson, gives a monopole density $\rho_M(\beta)$ that decays exponentially over five decades from 0.237 at $\beta = 0.05$ down to 5×10^{-5} at $\beta = 1.0$. A Polyakov dilute-gas fit $\rho_M(\beta) = A e^{-c\beta}$ returns $c = 9.40$ at $n_{\max} = 2$ and $c = 8.95$ at $n_{\max} = 3$ (both $R^2 = 0.981$, exponent $n_{\max}$ -independent within 5%), confirming a 3D dilute-plasma signature rather than a 4D condensation transition. (vi) *Full Monte Carlo Wilson loops* (Sprint XCWG-G, May 2026, NUANCED WEAK PASS): full-MC sampling of the compact $U(1)$ Wilson measure gives ensemble-mean string tension $\sigma_{\text{ens}}(\beta) > 0$ monotone decreasing at every $\beta \in [0.1, 10]$ on Rule B at $n_{\max} \in \{2, 3\}$, qualitatively consistent with 3D compact $U(1)$; *but* a joint area-plus-perimeter fit shows that the genuine area-law coefficient σ_{comb} is statistically zero, and the apparent confinement signal is carried by the perimeter coefficient μ_{comb} , not the area coefficient – a finite-volume effect on the irregular non-cubic Rule B graph at the accessible cutoffs. The Polyakov cross-check $\sigma \propto \exp(-c_\rho/2)\beta$ from XCWG-F passes an exponential-shape test ($R^2 = 0.83\text{--}0.88$) but with

rate constant $c_\sigma^{\text{MC}} = 1.20\text{--}1.63$, a factor of $3\text{--}4\times$ smaller than the predicted $c_\rho/2 = 4.70$, attributable to the non-cubic dual-lattice topology of Rule B (mean ~ 4.4 plaquettes per vertex) preventing the textbook continuum Polyakov derivation from applying verbatim. Clean area-vs-perimeter separation requires $n_{\text{max}} \geq 4$ ($V = 60$). (vii) *Finite-temperature Polyakov-loop analog* (Sprint XCWG-H, May 2026, CLEAN PASS): a brand-new product-graph construction $G_B \times C_{N_t}$ – the discrete analog of $\mathbb{R}^3 \times S_\beta^1$ – supplies a Euclidean-time-compactified substrate on which the Polyakov loop is well-defined. Homologically, the construction at $N_t \in \{4, 8\}$ has $\dim \ker(d_1) = 1$, with the one extra cocycle being *exactly the Polyakov-loop class*: the temporal-winding holonomy survives Gauss-law gauge equivalence as a homologically forced observable, not by hand-pick. The $N_t = 2$ case is correctly excluded as a kinematic degeneracy ($P \equiv 1$ by single-edge backtracking, $\dim \ker(d_1) = 0$). Monte Carlo at $(\beta, N_t) \in \{0.3, 1.0, 3.0\} \times \{4, 8\}$ returns $|\langle P \rangle|$ consistent with 0 within 2σ at all six non-degenerate cells, and the two-point Polyakov correlator decays approximately exponentially with spatial graph distance, with the correct N_t -dependence (per-step decay ratio ~ 0.7 at $N_t = 8$ vs ~ 0.8 at $N_t = 4$ at $\beta = 3$, consistent with hotter systems having longer spatial correlation lengths). This is the canonical signature of *persistent confinement at finite temperature*: the would-be Kaluza-Klein reduction $S_\beta^1 \rightarrow$ effective 2D theory is also confined (Wilson 1974), so no deconfinement transition occurs as T rises, exactly the 3D compact $U(1)$ behaviour that distinguishes it from the 4D Svetitsky-Yaffe finite- T deconfinement. The seven witnesses jointly establish structural compatibility with 3D compact $U(1)$ at leading order, non-perturbatively at the dilute-monopole-plasma level, and at finite temperature via persistent confinement; six are clean passes, the sixth (XCWG-G) is a nuanced weak pass with explicit finite-volume scope. Honest scope: $n_{\text{max}} \in \{2, 3, 4\}$ and finite- β Monte Carlo give sprint-level resolution, not thermodynamic-limit physics. Two v4 caveats are sharpened in v5 by completed sprints XCWG-G (extended to $n_{\text{max}} = 4$) and XCWG-J (Polyakov-rate refinement on Rule B’s dual graph G_B^*). The XCWG-G $n_{\text{max}} = 4$ extension reveals that σ_{comb} in the joint $\sigma A + \mu L$ fit does *not* clean up at the larger graph – it sharpens to $z = -48$ via a non-monotonic $A \leftrightarrow L_{\text{min}}$ coupling driven by cluster topology, not gauge dynamics; the cleanest fixed- $L = 6$ datum gives $\sigma = -0.004 \pm 0.003$ (consistent with zero), the honest area-law coefficient at the available statistics. The XCWG-J refinement resolves the apparent “ $3\text{--}4\times$ disagreement” as a category mismatch between observables: numerically, $c_{\sigma_{\text{ens}}} = 1.196$ and $c_{\mu_{\text{comb}}} = 1.242$ agree within 4%, so XCWG-G measures the perimeter-coefficient rate, not the Polyakov-predicted string-tension rate; the latter is $\sigma \sim 5 \times 10^{-9}$ at $\beta = 1.5$ (four orders below the MC detection floor), consistent with the $\sigma_{\text{comb}} \approx 0$ observation. Non-cubic dual-lattice corrections (dual graph average degree 9.60 vs $\mathbb{Z}^3$ ’s 6.0, spectral gap $\lambda_2 = 2.138$ vs 3.0) enter the Polyakov prefactor σ_0 , not the rate constant $c_\sigma = c_\rho/2$. The seven-witness reading remains “structurally compatible with 3D compact $U(1)$ including its non-perturbative monopole-plasma mechanism and finite-temperature persistent-confinement signature,” not “uniquely identified with 3D up to universality.”

1 Introduction

The GeoVac framework discretises Coulomb bound-state quantum mechanics onto a finite graph whose continuum limit is the unit three-sphere S^3 [1]. A series of structural readings has placed this graph inside the Wilson lattice gauge vocabulary: Paper 25 [5] constructed $U(1)$ Wilson gauge on the scalar Fock graph with link variables on the radial $T_\pm$ and angular $L_\pm$ ladder edges; Paper 30 [8] promoted the construction to non-abelian $SU(2)$ via the coincidence $S^3 = SU(2)$; Sprint ST-SU3 (May 2026, `debug/st_su3_wilson_memo.md`) constructed Wilson $SU(3)$ on the Bargmann–Segal S^5 graph of Paper 24 [4], with the universal weak-coupling kinetic coefficient $1/(4N_c)$. The three constructions exhibit all Standard Model compact gauge groups as Wilson constructions on natural sub-manifolds.

A subsequent RG-flow diagnostic (Sprint RG-A, May 2026) identified a structural gap in the

Paper 25 construction. The scalar Fock graph factorises as a disjoint union of 2D rectangular grids

$$G_{\text{scalar}} = \bigsqcup_{\ell=0}^{n_{\text{max}}-1} P_{n_{\text{max}}-\ell} \times P_{2\ell+1}, \quad (1)$$

one per orbital quantum number ℓ , because the scalar adjacency rules ($T_{\pm}$ raise/lower n at fixed (ℓ, m) ; $L_{\pm}$ raise/lower m at fixed (n, ℓ)) never cross ℓ -shells. Migdal–Kadanoff block-spin on the scalar graph therefore reproduces Polyakov’s 2D compact $U(1)$ flow per-block, with $\beta_{\text{eff}} \rightarrow 0$ monotonically, and no genuine 3D phase content emerges. The structural obstruction is geometric: the scalar Wilson construction has spectral dimension $d_s \approx 2$ per ℓ -block, well below the $d_s = 3$ of S^3 .

To recover 3-dimensionality, the graph needs cross- ℓ hopping edges. Paper 29 §RH-C [7] supplies exactly such a substrate: the Dirac graph “Rule B” with vertices labeled by Dirac quantum numbers $v = (n_{\text{fock}}, \kappa, m_j)$ and edges given by the standard atomic E1 dipole selection rule

$$v \sim_B v' \iff |\Delta n_{\text{fock}}| \leq 1, |\Delta \ell| = 1, |\Delta j| \leq 1, |\Delta m_j| \leq 1, v \neq v'. \quad (2)$$

The strict $\Delta \ell = \pm 1$ requirement forces *every* Rule B edge to flip ℓ -parity. Within-shell edges are forbidden by construction; the graph mixes all ℓ -shells into a single connected component.

This paper specifies $U(1)$ Wilson lattice gauge theory on Rule B, places it in the GeoVac Wilson program of Papers 25 / 30 / Sprint ST-SU3, and tests structural compatibility with 3D continuum compact $U(1)$ – the Polyakov 1977 / Banks–Myerson–Kogut 1977 universality class characterised by permanent confinement at every $\beta > 0$, driven by a magnetic-monopole plasma whose density and induced string tension decrease exponentially with β but never vanish at finite coupling, and by the absence of any finite-temperature deconfinement transition (the would-be Kaluza-Klein reduction $S^1_{\beta} \rightarrow 2\text{D}$ is also confined). The question we test is whether the Rule B construction is structurally compatible with that universality class, in the sense that seven standard structural witnesses (effective dimension, RG flow direction, Gaussian-sector Wilson loop, strong-coupling area law, non-perturbative Polyakov monopole density, full-MC Wilson loops, and finite-temperature Polyakov loops) each return values consistent with the 3D-compact- $U(1)$ prediction and inconsistent with the 4D phase-transition prediction.

The verdict, stated up front: **structural compatibility holds across all seven witnesses, with six clean passes and one nuanced weak-pass**. The first four are leading-order or coarse-approximation tests; the fifth (the non-perturbative monopole density via DeGrand–Toussaint Monte Carlo) is the canonical Polyakov-mechanism discriminator and lies at the non-perturbative level; the sixth (full-MC Wilson loops on the compact $U(1)$ measure) is nuanced – the ensemble-mean string tension is positive and monotone decreasing in β on Rule B at $n_{\text{max}} \leq 3$, but a joint area-plus-perimeter fit reveals that the apparent confinement signal is carried by the perimeter coefficient, not the area coefficient, on the available graph sizes; the seventh (Polyakov loops on a product graph $G_B \times C_{N_t}$) is a clean pass via the homological Polyakov class (the one extra first-cohomology cocycle not generated by plaquettes at $N_t \geq 3$) and the MC verification that $\langle P \rangle = 0$ within 2σ at all non-degenerate (β, N_t) . The 3D-vs-4D distinction in compact lattice gauge theory is fundamentally non-perturbative: it lives in the existence of a deconfined phase at finite β_c in 4D vs the permanent-confinement plasma at every $\beta > 0$ in 3D, and at finite temperature in the absence of a Svetitsky-Yaffe deconfinement transition in 3D vs its presence in 4D. Leading-order character expansion and Migdal–Kadanoff RG cannot resolve this; the monopole-density, full-MC-Wilson-loop, and Polyakov-loop observables can. The seven-witness reading adds three structurally distinct probes on top of the four leading-order witnesses, each cleanly compatible with the 3D-compact- $U(1)$ prediction at the level it accesses.

The conclusion is honestly bounded in six ways. (i) The Monte Carlo verification samples $n_{\max} \in \{2, 3\}$, not the thermodynamic limit; size scaling is consistency-checked across the two cutoffs (XCWG-F decay exponent c stable within 5%; XCWG-G σ_{ens} qualitatively preserved; XCWG-H verdict identical at $N_t = 4$ and $N_t = 8$) but not asymptotic. (ii) The XCWG-F MC detection floor caps the directly measured range at $\beta \leq 1.0$; the conclusion that $\rho_M > 0$ on $[1.5, \infty)$ rests on exponential extrapolation of a fit with $R^2 = 0.981$ over five decades. (iii) Per CLAUDE.md §1.5, the discrete graph and the continuum manifold are dual descriptions connected by the proven Fock projection; this paper does not assert ontological priority of either. (iv) XCWG-G is a *nuanced weak-pass*: the full-MC ensemble-mean $\sigma_{\text{ens}}(\beta) > 0$ at every β is structurally the right 3D-compact- $U(1)$ shape, but the joint area-plus-perimeter fit identifies the available graph sizes as perimeter-dominated; a clean separation of area-law from perimeter-law content requires $n_{\max} \geq 4$. (v) The XCWG-G Polyakov-rate cross-check quantitatively disagrees with XCWG-F: the measured exponential rate constant $c_\sigma^{\text{MC}} = 1.20\text{--}1.63$ is $3\text{--}4\times$ smaller than XCWG-F’s predicted $c_\rho/2 = 4.70$, attributable to Rule B’s non-cubic dual-lattice topology (~ 4.4 plaquettes per vertex versus ~ 1 on $\mathbb{Z}^4$) preventing the textbook Polyakov continuum derivation from applying verbatim. The qualitative exponential-shape agreement holds; the quantitative rate disagreement is recorded as an honest structural feature. (vi) XCWG-H’s finite-temperature verdict is at $n_{\max} = 2$ only and uses a first-pass Metropolis MC; the spatial graph diameter is 4, which limits the dynamic range of the two-point Polyakov correlator and prevents a clean string-tension extraction at this size. What the seven witnesses jointly verify is *absence of the 4D phase transition* (witnesses 2 and 4), *presence of the 3D area-law mechanism at leading order* (witness 4), *presence of the 3D Polyakov dilute-plasma signature in the monopole density* (witness 5), *qualitative compatibility of the full compact Wilson measure with 3D string-tension behaviour subject to finite-volume perimeter-dominance* (witness 6), and *persistent confinement at finite temperature with the homological Polyakov-class structure* (witness 7), not a unique identification with 3D up to universality.

The structural content is concrete, scales cleanly, and slots into a definite position in the GeoVac Wilson program: $U(1)$ on Hopf S^3 (Paper 25, per- ℓ 2D) $\rightarrow$ $SU(2)$ on $S^3 = SU(2)$ (Paper 30) $\rightarrow$ $SU(3)$ on S^5 Bargmann (Sprint ST-SU3) $\rightarrow$ $U(1)$ on Dirac Rule B with cross- ℓ structure (this paper, 3D-compatible).

2 The Dirac-Rule-B graph

2.1 Vertex and edge data

The vertices of the Rule B graph at cutoff $n_{\max}$ are Dirac orbital labels $(n_{\text{fock}}, \kappa, m_j)$ with $n_{\text{fock}} \in \{1, 2, \dots, n_{\max}\}$, $\kappa \in \mathbb{Z} \setminus \{0\}$ encoding the orbital quantum number ℓ and total angular momentum j via

$$\ell(\kappa) = \begin{cases} -\kappa - 1 & \kappa < 0 \\ \kappa & \kappa > 0 \end{cases}, \quad j = |\kappa| - \frac{1}{2}, \quad (3)$$

and $m_j \in \{-j, -j+1, \dots, j\}$ a half-integer with $|m_j| \leq j$. The vertex count per shell is $2n^2$ (the standard Dirac multiplicity); total vertex count is $V(n_{\max}) = \sum_{n=1}^{n_{\max}} 2n^2$.

Edges are defined by (2), reproducing the canonical atomic E1 dipole selection rule. Condition $|\Delta\ell| = 1$ forces every edge to cross ℓ -shells.

2.2 Graph properties

Numerical graph properties at the cutoffs computed in this paper:

$n_{\max}$	V	E	c	β_1	diameter	girth	$\langle q_e \rangle$	$L = 4$ plaquettes
2	10	20	1	11	4	4	8.80	44
3	28	106	1	79	6	4	37.51	994
4	60	312	1	253	8	4	–	5,047
5	110	692	1	583	10	4	–	14,881

Here c is the number of connected components, $\beta_1 = E - V + c$ is the first Betti number (the dimension of the cycle space), and $\langle q_e \rangle$ is the mean number of length-4 plaquettes incident to each edge. The Rule B graph is connected at every tested $n_{\max}$; this is the most consequential single property, distinguishing it from the scalar Fock graph whose component count equals $n_{\max}$ (one per ℓ -shell).

2.3 Comparison to the scalar Fock graph

Side-by-side at matched cutoff (scalar Fock data from Papers 25, 29):

Property	Scalar $n_{\max} = 3$	Rule B $n_{\max} = 3$	Scalar $n_{\max} = 4$	Rule B $n_{\max} = 4$
V	14	28	30	60
E	13	106	34	312
Components c	3	1	4	1
β_1	2	79	8	253
Cross- ℓ edge fraction	0%	100%	0%	100%
Cross-shell plaquette fraction	0%	100%	0%	100%
$L = 4$ plaquettes	2	994	8	5,047
Plaquettes/edge ratio	0.15	9.38	0.24	16.18

Three structural facts hold uniformly across $n_{\max} \in \{2, 3, 4\}$.

Connectivity. The scalar Fock graph splits as $\bigoplus_{\ell} P_{n_{\max}-\ell} \times P_{2\ell+1}$, one disconnected 2D rectangular grid per ℓ . Rule B is connected because every dipole edge links neighbouring ℓ -shells.

Cross- ℓ structure. On the scalar graph, 0% of edges cross ℓ -shells; on Rule B, 100% of edges cross ℓ -shells. The fraction is *categorical*, not statistical: it follows from the selection rules of the respective adjacency definitions.

Plaquette density. The plaquette/edge ratio measures how many 2-cells each edge bounds. Reference values: 2D square lattice 0.5, 3D cubic lattice 1.5, 4D hypercubic 3.0. Rule B at $n_{\max} = 3$ has ratio 9.38, well above the 4D hypercubic value; the scalar graph at the same cutoff has ratio 0.15.

2.4 Plaquette signatures

At $n_{\max} = 3$, every length-4 plaquette in Rule B is bipartite under ℓ -parity (every edge flips parity, so 4-cycles have alternating parity). The 994 plaquettes partition by sorted ℓ -multiset into three classes:

ℓ -signature	Count	Fraction	Type
(0, 0, 1, 1)	272	27.4%	two-shell, $s \leftrightarrow p$ tetrahedron
(0, 1, 1, 2)	542	54.5%	three-shell, “rotating-shell”
(1, 1, 2, 2)	180	18.1%	two-shell, $p \leftrightarrow d$ tetrahedron

The dominant plaquette signature spans three ℓ -shells. This is the combinatorial witness for cross-shell content beyond what the scalar construction can supply: 54.5% of Rule B’s 4-plaquettes at $n_{\max} = 3$ are s - p - d cycles, structurally distinct from any 2D within-shell plaquette.

2.5 Bipartiteness and girth

Rule B is bipartite under the $\mathbb{Z}_2$ coloring by $(-1)^\ell$: every edge flips ℓ -parity, so every closed walk has even length. The shortest cycle is length 4 at every tested $n_{\max}$.

This bipartiteness is structurally important. It implies (i) no triangle plaquettes, so all 2-cells are quadrilaterals; (ii) under the $\mathbb{Z}_2$ vertex-coloring, the Wilson partition function on Rule B has a natural staggered structure – but no spinor-doubling artifact, because all Rule B objects are simplicial in the linear algebra (signed incidence, edge Laplacian, plaquette boundary) rather than fermion-spinor.

2.6 Spectral genealogy

Paper 29 §RH-C established that the Ihara zeta of Rule B is strictly sub-Ramanujan at $n_{\max} = 3$ (deviation -0.119 from the Kotani–Sunada Ramanujan bound) and supra-Ramanujan at $n_{\max} = 4$ (deviation $+1.53$); the Dirac-Rule-B Ihara zeta factors as $(s \pm 1)^{79} \cdot P_{22}(s^2) \cdot P_{24}(s^2)$ at $n_{\max} = 3$, a $22 + 24$ block decomposition under the $m_j \rightarrow -m_j$ $\mathbb{Z}_2$ graph automorphism, exactly the Paper 29 Hypothesis 1 prediction.

3 The Wilson $U(1)$ construction on Rule B

3.1 Link variables and gauge transformation

Each undirected edge $\{v, v'\} \in E$ is promoted to two oriented edges $e = v \rightarrow v'$ and $\bar{e} = v' \rightarrow v$, with the canonical choice $v \rightarrow v'$ for the smaller node index as source. A $U(1)$ *link variable* is a phase

$$U_e = e^{i\theta_e} \in U(1), \quad U_{\bar{e}} = U_e^{-1}, \quad \theta_e \in [0, 2\pi). \quad (4)$$

The free configuration space of the gauge theory is $U(1)^{|E|}$, with Haar measure $\prod_e d\theta_e/(2\pi)$.

A gauge transformation is a function $\chi : V \rightarrow \mathbb{R}/2\pi\mathbb{Z}$, $v \mapsto \chi_v$, acting on link variables by

$$\theta_e \mapsto \theta_e + (\chi_{v'} - \chi_v) \pmod{2\pi} \quad \text{for } e = v \rightarrow v', \quad (5)$$

and on matter fields $\psi_v : V \rightarrow \mathbb{C}$ by $\psi_v \mapsto e^{i\chi_v} \psi_v$. Gauge-invariant observables are products of link variables around closed walks (holonomies); the basic one is the plaquette holonomy $U_P = \prod_{e \in P} U_e$.

A note on the spinor labels. The vertices of Rule B carry Dirac quantum numbers (κ, m_j) , but in this paper the spinor structure plays the role of an indexing convention: a matter field ψ_v is a complex scalar attached to the vertex, not a 4-component Dirac spinor with internal indices. The $U(1)$ acts by an overall phase at each orbital. The natural larger gauge group $U(1) \times \text{SU}(2)$ (charge times intrinsic spin rotation) is out of scope for the abelian construction studied here. The companion non-abelian construction is Paper 30’s $\text{SU}(2)$ Wilson on the scalar S^3 graph.

3.2 Wilson action

The Wilson action on Rule B is

$$S_W[\theta] = \beta \sum_{P \in \mathcal{P}} (1 - \cos \theta_P), \quad \theta_P = \sum_{e \in P} \theta_e \quad (6)$$

where $\mathcal{P}$ is the set of primitive non-backtracking closed 4-walks on G_B , and θ_P is the oriented sum of link phases around the plaquette. The action is real, non-negative, and bounded below by 0 at the trivial vacuum $\theta_e \equiv 0$. It is gauge-invariant because θ_P is a sum of $(\chi_{v'} - \chi_v)$ -differences around a closed loop and these telescope.

For $\beta \rightarrow 0$ (strong coupling) the path integral is dominated by random link configurations; for $\beta \rightarrow \infty$ (weak coupling) it is dominated by configurations near the trivial vacuum.

3.3 Signed incidence and the two Laplacians

Choose the canonical orientation $v \rightarrow v'$ for each edge. The signed incidence matrix $B \in \mathbb{Z}^{V \times E}$ has entries

$$B_{v,e} = \begin{cases} -1 & v = \text{source}(e) \\ +1 & v = \text{target}(e) \\ 0 & \text{otherwise.} \end{cases} \quad (7)$$

The two Laplacians are

$$L_0 = BB^\top \in \mathbb{Z}^{V \times V}, \quad L_1 = B^\top B \in \mathbb{Z}^{E \times E}. \quad (8)$$

L_0 is the standard combinatorial graph Laplacian on 0-cochains (vertex fields); L_1 is the edge Laplacian on 1-cochains (link fields). Both are integer matrices, so all eigenvalues are algebraic over $\mathbb{Q}$.

The Hodge identity asserts that the nonzero spectra of L_0 and L_1 coincide with multiplicity:

$$\sigma(L_0) \setminus \{0\} = \sigma(L_1) \setminus \{0\}, \quad (9)$$

together with the kernel dimensions $\dim \ker L_0 = c = 1$ (connectedness) and $\dim \ker L_1 = \beta_1$ (cycle rank). We verified this numerically at $n_{\max} \in \{2, 3\}$ in the pilot driver (`debug/xcwg_u1_wilson_rule_b_pilot.py`):

- $n_{\max} = 2$: $\dim \ker L_0 = 1$, $\dim \ker L_1 = 11 = \beta_1$, nonzero spectra of length 9 agree to $< 10^{-9}$.
- $n_{\max} = 3$: $\dim \ker L_0 = 1$, $\dim \ker L_1 = 79 = \beta_1$, nonzero spectra of length 27 agree to $< 10^{-9}$.

The Hodge identity is structural (it follows from L_0 and L_1 sharing the SVD spectrum of B); the numerical check certifies the implementation rather than uncovering a new fact.

3.4 The plaquette boundary operator and the correct Wilson kinetic

A subtle point worth flagging explicitly: *the right Wilson kinetic operator on the gauge-fixed subspace is not L_1 but the co-exact part of the Hodge Laplacian*. The edge Laplacian $L_1 = B^\top B$ has kernel containing every closed cycle ($BC = 0$ for any 1-cycle C), so when applied to Wilson-loop observables it produces a contact-zero artifact. The right object is the *co-exact* part

$$K = d_1^\top d_1, \quad (10)$$

where $d_1 : \mathbb{Z}^{|E|} \rightarrow \mathbb{Z}^{|P|}$ is the plaquette boundary operator with $(d_1)_{P,e} = +1$ if e appears in P in the canonical orientation, -1 if reversed, 0 otherwise.

On a graph equipped with a 2-complex from its primitive 4-cycles, the Hodge–Helmholtz decomposition reads

$$L_1 = d_0^\top d_0 + d_1^\top d_1 \quad (\text{with } d_0 = B^\top), \quad (11)$$

splitting edge cochains into the longitudinal (image of d_0 , pure-gauge) and transverse (image of $d_1^\top$, physical photon) sectors. The Faddeev–Popov gauge fixing removes the longitudinal modes, leaving $K = d_1^\top d_1$ as the quadratic kinetic operator. We verified at $n_{\max} \in \{2, 3, 4, 5\}$ that $\text{rank}(K) = \text{rank}(d_1) = \beta_1$, matching the Hodge theorem on the 2-complex.

This convention will matter explicitly when computing Gaussian Wilson loops in §6.1; using L_1^+ instead of K^+ would give contact zeros that mask the actual scaling.

3.5 Quadratic expansion and the weak-coupling kinetic term

Linearising the Wilson action (6) around the trivial vacuum $\theta_e = 0$:

$$1 - \cos \theta_P = \frac{1}{2} \theta_P^2 + O(\theta_P^4), \quad (12)$$

so

$$S_W[\theta] = \frac{\beta}{2} \sum_P \theta_P^2 + O(\theta^4) = \frac{\beta}{2} \theta^\top K \theta + O(\theta^4), \quad (13)$$

with $K = d_1^\top d_1$ from (10). On the gauge-fixed subspace $\theta \perp \text{im}(d_0)$, K acts as L_1 ; this is the abelian sibling of Paper 30’s SU(2) Theorem 2.

3.6 Plaquette counts and the structural difference from Paper 25

The headline structural difference between the Rule B Wilson construction and Paper 25’s scalar Wilson construction:

Property	Paper 25 scalar (Hopf)	Rule B (Dirac dipole)
V at $n_{\max} = 3$	14	28
E at $n_{\max} = 3$	13	106
c at $n_{\max} = 3$	3	1
β_1 at $n_{\max} = 3$	2	79
$L = 4$ plaquette count at $n_{\max} = 3$	2	994
Smallest plaquette length	6	4
Block decomposition of L_1	per ℓ -shell	κ -mixing, no block structure

Two of these are operationally consequential.

Length-4 plaquettes appear at $n_{\max} = 2$. Paper 25’s scalar graph has no length-4 plaquettes at any tested cutoff – its smallest plaquette is length 6 (the hexagonal $\ell = 1$ loop at $n_{\max} = 3$). Rule B has 44 length-4 plaquettes already at $n_{\max} = 2$ (where Paper 25’s graph is a forest, $\beta_1 = 0$). This is the discrete-geometric reason to expect Rule B to host non-trivial RG flow that the scalar graph cannot.

No per- ℓ -block decomposition. Paper 25’s L_1 is block-diagonal by ℓ -shell because every edge preserves ℓ . Rule B’s L_1 has no such block structure because every edge flips ℓ . The Wilson kinetic operator $K = d_1^\top d_1$ is therefore intrinsically cross- ℓ on Rule B, with no 2D-grid reduction available.

4 Spectral dimension

The most direct structural witness for 3-dimensionality is the spectral dimension d_s of the graph Laplacian, extracted from the heat-kernel return probability.

4.1 Method

For graph Laplacian $L = L_0$ with non-negative eigenvalues $\{\lambda_n\}_{n=1}^V$, the heat-kernel trace is $K(t) = \sum_n e^{-\lambda_n t} = \text{Tr } e^{-Lt}$, and the return probability per node is $P(t) = K(t)/V$. For a continuous Laplacian on a d -dimensional Riemannian manifold, $P(t) \sim Ct^{-d/2}$ at small t , so

$$d_s = -2 \frac{d \log[P(t) - P(\infty)]}{d \log t}. \quad (14)$$

The constant-mode contribution $P(\infty) = c/V$ is subtracted before the log-log slope; we fit in the intermediate window where $P(t) - P(\infty)$ lies between $5P(\infty)$ and 0.5 .

A Weyl-law cross-check fits $N(\lambda) \sim C' \lambda^{d/2}$ over the middle 50% of the spectrum, yielding d_W . A diffusion cross-check fits $\langle r^2(t) \rangle_{v_0} \sim t^{2/d_w}$, with $d_w = 2$ for ordinary diffusion.

4.2 Numerical results

Heat-kernel d_s , Weyl d_W , and diffusion slopes at $n_{\max} \in \{3, 4, 5\}$ for both Rule B and the scalar Fock graph:

Graph	$n_{\max}$	V	c	d_s	d_W	$\langle r^2 \rangle$ slope
Rule B	3	28	1	1.86	2.48	0.94
Rule B	4	60	1	2.27	3.46	0.91
Rule B	5	110	1	2.54	3.72	0.87
Scalar	3	14	3	1.68	0.98	0.98
Scalar	4	30	4	1.76	1.60	1.01
Scalar	5	55	5	1.79	1.71	1.03

4.3 Reading

Three structural facts.

Rule B's d_s climbs monotonically toward 3. The increments $\Delta d_s \approx 0.3$ per $n_{\max}$ step are roughly linear, and linear extrapolation suggests $d_s \approx 3$ near $n_{\max} \approx 7$ –8. Whether this is sustained or slowed by saturation effects is a finite-size question; the *direction* of the trend is robust.

The Weyl law crosses 3 between $n_{\max} = 4$ and $n_{\max} = 5$. $d_W = 2.48 \rightarrow 3.46 \rightarrow 3.72$ shows the eigenvalue density growing like $N(\lambda) \sim \lambda^{3/2}$ at intermediate λ , which is the S^3 Weyl growth. The heat-kernel and Weyl-law estimators bracket the continuum value $d = 3$ at $n_{\max} = 5$: heat-kernel below, Weyl above. Their discrepancy (≈ 1.2 at $n_{\max} = 5$) is a known finite-size phenomenon related to which window of the spectrum each estimator weights; both indicators *agree on the trend* that Rule B is climbing toward $d = 3$.

The scalar Fock graph plateaus at $d_s \approx 1.8$. The scalar graph sits at $d_s = 1.68 \rightarrow 1.79$, well below 2, and does *not* move toward 3 with $n_{\max}$. The gap is structural: the scalar graph decomposes into disjoint 2D grids, and any union of finite 2D grids has $d_s \rightarrow 2$ from below. This is consistent with the structural argument that the scalar Fock graph cannot host genuine 3-dimensionality.

Diffusion is approximately normal on both graphs. The diffusion slopes $\langle r^2 \rangle$ are close to 1, consistent with walk dimension $d_w \approx 2$. Rule B is slightly subdiffusive (slope 0.87 at $n_{\max} = 5$), which via the Einstein relation $d_s = 2d_f/d_w$ is consistent with $d_s < d_w$.

The first witness for 3D-compatibility passes: **Rule B's effective dimension is between 2 and 3 at finite $n_{\max}$, monotonically trending toward 3, while the scalar Fock graph is structurally trapped near 2.**

5 Migdal–Kadanoff RG flow

The second witness for 3D-compatibility is the direction of the Migdal–Kadanoff (MK) bond-moving block-spin renormalisation group flow. On a D -dimensional regular hypercubic lattice with block size b , the canonical MK recursion for $U(1)$ is

$$\frac{I_1(\beta_{\text{eff}})}{I_0(\beta_{\text{eff}})} = \left[\frac{I_1(\beta)}{I_0(\beta)} \right]^k, \quad k = b^{D-1}, \quad (15)$$

where I_n are modified Bessel functions. The exponent k has a direct combinatorial interpretation: it counts the number of plaquettes sharing an integrated bond after bond-moving. At $b = 2$: 2D square gives $k = 2$ (Polyakov's plane-rotator confinement); 3D cubic gives $k = 4$ (Polyakov–BMK permanent confinement, though MK misses the deconfinement transition by being unable to resolve monopoles); 4D hypercubic gives $k = 8$.

5.1 Adaptation to an irregular graph

For an irregular graph like Rule B (degree spread 2–12 at $n_{\max} = 3$, plaquette-per-edge spread 6–62), bond-moving cannot be applied uniformly – there is no canonical parallel-transport direction. Instead we work with the *plaquette-per-edge incidence number* q_e directly: for each edge e , count the number of distinct length-4 plaquettes containing e , and take the mean

$$k_{\text{eff}} = \langle q_e \rangle_E = \frac{1}{|E|} \sum_{e \in E} q_e. \quad (16)$$

On a 2D square lattice every interior edge has $q_e = 2$, recovering $k = 2$; on 3D cubic, $q_e = 4$. The Rule B values are tabulated in §2.2: $k_{\text{eff}} = 8.80$ at $n_{\max} = 2$ and $k_{\text{eff}} = 37.51$ at $n_{\max} = 3$ – well above the cubic-lattice $k = 4$.

Three alternative choices ($k = q_e - 1$, geometric mean, $q_{\max}$) all give the same qualitative answer at both tested cutoffs.

5.2 Strong/weak-coupling asymptotics

The MK recursion reproduces standard limits in two regimes:

- Strong coupling ($\beta \rightarrow 0$): $\beta_{\text{eff}} \approx 2^{1-k} \beta^k$, verified to $\leq 0.25\%$ at $\beta = 0.03$ for $k = 8.8$.
- Weak coupling ($\beta \rightarrow \infty$): $\beta_{\text{eff}} \approx \beta/k$, verified to $\leq 2\%$ at $\beta = 50$ for $k = 8.8$.

5.3 Flow direction and absence of non-trivial fixed point

Cutoff	k_{eff}	Non-trivial fixed point?	Iter. to $\beta < 10^{-10}$ from $\beta_0 = 10$
Rule B $n_{\text{max}} = 2$	8.80	no	≤ 3
Rule B $n_{\text{max}} = 3$	37.51	no	≤ 2
2D scalar (textbook)	2.0	no (Polyakov)	~ 30
3D cubic (textbook MK)	4.0	no (MK misses 3D deconf.)	~ 10

At both Rule B cutoffs, $g(\beta) := \beta_{\text{eff}}(\beta) - \beta < 0$ on $\beta \in [10^{-3}, 10^2]$ across a 2001-point log-spaced grid, with zero sign changes. The recursion has $\beta = 0$ as the unique attractive fixed point and $\beta = \infty$ as an unstable one. There is no non-trivial fixed point.

5.4 Reading: consistent with permanent confinement

The MK flow direction matches the Polyakov–Banks–Myerson–Kogut prediction for 3D compact $U(1)$: monotone flow to $\beta = 0$, no UV fixed point, permanent confinement at every $\beta > 0$ [13, 14, 15].

Honest caveats. (i) MK is universal in a coarse sense: every $U(1)$ MK flow on a plaquette graph with $k_{\text{eff}} > 1$ goes monotonically to $\beta = 0$. The qualitative direction is therefore identical for 2D scalar, 3D cubic, and Rule B at the tested cutoffs. The 3D-vs-4D distinction is invisible to MK because MK cannot resolve monopoles (Polyakov 1977’s mechanism is non-perturbative in the monopole density).

(ii) What *is* discriminating in the MK data is the *absence* of a non-trivial fixed point at the tested cutoffs. The 4D compact $U(1)$ MK recursion (with $k = 8$ on the hypercubic lattice) also flows to $\beta = 0$ qualitatively; the real 4D phase transition at $\beta_c \approx 1.01$ is a non-perturbative effect MK is structurally unable to resolve. Rule B’s MK flow shape is therefore consistent with both 3D and 4D compact $U(1)$ – it does not by itself uniquely identify the universality class.

(iii) The Rule B flow is quantitatively much faster than the textbook cases: at $n_{\text{max}} = 3$, $k_{\text{eff}} = 37.5$ produces strong-coupling recursion $\beta_{\text{eff}} \propto \beta^{37.5}$, an extraordinarily steep flow. This reflects the high plaquette density of Rule B (plaquette/edge ratio 9.4 vs cubic-lattice 1.5), not a distinct phase structure.

The second witness for 3D-compatibility passes: **the MK flow on Rule B points to permanent confinement at every β , with $\beta = 0$ as the unique attractive fixed point.** This is consistent with the Polyakov–BMK prediction for 3D compact $U(1)$ and inconsistent with a 4D phase-transition scenario at finite β_c . The reading is structural compatibility, not unique identification.

6 Wilson loop diagnostics

The remaining two structural witnesses are observables of Wilson loops themselves – the canonical confinement order parameter in lattice gauge theory.

6.1 Gaussian (free-photon) Wilson loops

At weak coupling ($\beta \rightarrow \infty$), Wilson loops factorize through the Gaussian quadratic form

$$\langle W(C) \rangle \sim \exp\left(-\frac{1}{2\beta} C^\top K^+ C\right), \quad (17)$$

where C is the signed edge-indicator vector of the loop and K^+ is the Moore–Penrose pseudoinverse of $K = d_1^\top d_1$ from (10). The exponent $S(C) := C^\top K^+ C$ is the Gaussian Wilson loop action; its scaling with loop length L is the relevant order parameter in the free-photon sector.

We computed $\langle S(L) \rangle$ as a function of length $L \in \{4, 6, 8\}$ at $n_{\max} \in \{2, 3, 4\}$, and at $L \in \{4, 6\}$ for $n_{\max} = 5$. The enumeration is primitive non-backtracking closed walks under cyclic and reversal canonicalization; for each, $S(C)$ is computed from K^+ directly.

Methodology corrections. Three methodological issues were caught and corrected during this sprint: (i) a closure-step backtracking bug in the pilot walk enumeration ($v_{L-2} \neq v_0$ constraint missed, over-counting $L = 6$ walks by $\sim 4\times$ at $n_{\max} = 2$, fixed); (ii) a DFS start-vertex bias in capped enumerations (low-index vertices over-sampled, producing 55% running-mean drift at $n_{\max} = 4, L = 6$, fixed by randomised DFS order with fixed seed); (iii) a sub-sample-mean bias in the pilot $\langle S(L = 6) \rangle$ (deterministic ordering placed high- S walks first; fixed by full-population mean).

Numerical results (after corrections):

$n_{\max}$	$\langle S(L = 4) \rangle$	$\langle S(L = 6) \rangle$	$\langle S(L = 8) \rangle$	Scaling α
2	0.2477	0.3843	0.5687	1.18 ± 0.08
3	0.0795	0.1202	0.1618	1.03 ± 0.004
4	0.0501	0.0729	0.0929	0.89 ± 0.02
5	0.0392	0.0585	–	0.99 (2-point)

The scaling exponent α is the fit of $\langle S(L) \rangle \sim L^\alpha$ at fixed $n_{\max}$. Across all four cutoffs, $\alpha \in [0.89, 1.18]$ with mean $\bar{\alpha} \approx 1.0$. The fit quality is high ($R^2 \in [0.94, 1.00]$), the band is tight, and the direct ratio cross-check confirms perimeter-law predictions to $\leq 5\%$ ($S(6)/S(4) \rightarrow 1.5$, $S(8)/S(6) \rightarrow 1.33$).

Reading: 3D free-photon perimeter law. The result $\alpha \approx 1$ is the canonical signature of the *free-photon* sector in 3D abelian lattice gauge theory: in D dimensions the Coulomb-gauge propagator K^+ falls as $|x - y|^{-(D-2)}$, and a Wilson loop of length L has Gaussian action scaling as $L^{(D-2)} \cdot L = L^{D-1}$ for $D \leq 4$, modulo logarithmic corrections. In $D = 3$ this gives perimeter scaling $\alpha = 1$.

This sector *does not see compactness*. The Gaussian Wilson loop is computed by expanding $1 - \cos \theta_P \rightarrow \theta_P^2/2$, losing the 2π periodicity of the compact integration domain. The monopole excitations responsible for the compact- $U(1)$ area law (Polyakov 1977) are by construction absent from the Gaussian sector. The $\alpha \approx 1$ result therefore *confirms* that Rule B’s free-photon content is consistent with 3D geometry; it does *not* contradict permanent confinement of the full compact theory.

The earlier reading of the same numerical data (in the sprint pilot memo at `debug/xcwg.wilson.loop.scaling.m`) interpreted the $\alpha \approx 1$ result as “the second witness fails the area-law test, and the two-witness verification is therefore not established.” That reading conflated the free-photon sector (where perimeter

law is correct in 3D) with the full compact theory (where the area law is generated by monopoles). The correct discrimination is between 3D free-photon perimeter law and 4D free-photon L^2 scaling; the measured $\alpha \approx 1$ unambiguously selects the 3D case.

The third witness for 3D-compatibility passes: **the Gaussian Wilson-loop scaling on Rule B matches the 3D free-photon perimeter law $\alpha \approx 1$, with the area-law mechanism of compact $U(1)$ left to the strong-coupling expansion of §6.2.**

6.2 Strong-coupling character expansion

The proper compact- $U(1)$ confinement test is the leading-order strong-coupling character expansion. Expanding the Boltzmann factor as

$$e^{\beta \cos \theta_P} = \sum_{n \in \mathbb{Z}} I_n(\beta) e^{in\theta_P} \quad (18)$$

and integrating each link variable using $\int_{-\pi}^{\pi} e^{im\theta} d\theta / (2\pi) = \delta_{m,0}$, the leading nonzero contribution to a Wilson loop $W(C) = e^{i\theta_C}$ comes from tiling the loop with the minimum number $A(C)$ of plaquettes:

$$\langle W(C) \rangle_{\beta} = \left[\frac{I_1(\beta)}{I_0(\beta)} \right]^{A(C)} + O(\beta^{A(C)+1}). \quad (19)$$

The string tension is

$$\sigma(\beta) := -\ln(I_1(\beta)/I_0(\beta)), \quad (20)$$

giving the area law $\langle W(C) \rangle \sim e^{-\sigma(\beta)A(C)}$.

Asymptotic limits.

Regime	$I_1(\beta)/I_0(\beta)$	$\sigma(\beta)$
$\beta \rightarrow 0^+$	$\beta/2 + O(\beta^3)$	$\ln(2/\beta) \rightarrow +\infty$
$\beta \rightarrow \infty$	$1 - 1/(2\beta) + O(\beta^{-2})$	$1/(2\beta) \rightarrow 0^+$

Verification at high precision: $\sigma(0.01)_{\text{computed}} = 5.2983$ vs $\sigma_{\text{asympt}} = \ln(200) = 5.2983$ (relative error 2.4×10^{-6}); $\sigma(1000)_{\text{computed}} = 5.0025 \times 10^{-4}$ vs $\sigma_{\text{asympt}} = 5.0000 \times 10^{-4}$ (relative error 5×10^{-4}).

Construction by plaquette gluing. Computing $A(C)$ for an arbitrary loop is in general a non-trivial homology problem. We bypass this by constructing Wilson loops as oriented unions of plaquettes, so that the area is known by construction:

$$C(P_1, \dots, P_k) = \sum_{i=1}^k \epsilon_i v_{P_i}, \quad (21)$$

where $\epsilon_i \in \{\pm 1\}$ are orientation signs chosen so adjacent plaquettes cancel on shared edges. We verify C is closed ($BC = 0$) and simple (entries in $\{-1, 0, +1\}$), then evaluate (19) directly.

Numerical results. Cluster availability at $n_{\max} \in \{2, 3, 4\}$ and $A \in \{1, 2, 3, 4, 5\}$:

$n_{\max}$	Area A	Perimeters observed
2	1	$\{4\}$
2	2	$\{4, 6\}$
2	3	$\{4, 6, 8, 10\}$
2	4	$\{4, 6, 8, 10, 12\}$
3	1	$\{4\}$
3	2	$\{4, 6\}$
3	5	$\{6, 8\}$
4	5	$\{8, 10, 12, 14\}$

The string tension table is cutoff-independent at leading order:

β	I_1/I_0	$\sigma(\beta)$	$\langle W \rangle_{A=2}$
0.1	0.04994	2.9970	0.0025
0.3	0.14834	1.9083	0.0220
1.0	0.44639	0.8066	0.1993
3.0	0.80999	0.2107	0.6561
10	0.94860	0.05277	0.8999
30	0.98319	0.01695	0.9667
100	0.99499	0.00503	0.9900

Verdict: three structural facts.

- (i) *Area law holds at every tested $\beta > 0$. $\langle W(C) \rangle_\beta = (I_1/I_0)^{A(C)}$ for every constructed cluster C of area A , at every $n_{\max} \in \{2, 3, 4\}$.*
- (ii) *String tension is strictly positive at every $\beta > 0$. $\sigma(\beta) > 0$ uniformly, monotonically vanishing toward zero at large β but never reaching zero at finite coupling. This is permanent confinement, not phase-transitional.*
- (iii) *Asymptotic limits match the Polyakov–BMK prediction precisely. $\sigma \rightarrow \ln(2/\beta)$ at small β and $\sigma \rightarrow 1/(2\beta)$ at large β , both verified to relative error $\leq 5 \times 10^{-4}$.*

Perimeter-vs-area discriminator. The structural test distinguishing area law from perimeter law is whether Wilson loops of *equal area* but *different perimeter* give equal or different expectations. The leading-order character expansion predicts equal: $\langle W \rangle = (I_1/I_0)^A$ depends only on A , not on L . We verified at every cutoff and every area in the cluster table that the predicted $\langle W \rangle$ is identical across all observed perimeters (by construction at leading order; sub-leading corrections would introduce mild perimeter-dependent terms but would not affect the leading area-law verdict).

Honest caveat. The formula (19) is universal at leading order: it holds for any Wilson lattice gauge theory with the standard action, regardless of the underlying graph topology. What is verified is *structural compatibility* with the 3D-compact- $U(1)$ universality class: Rule B carries enough plaquette structure to admit the standard character expansion, and the resulting string tension has the correct asymptotic behavior. What is *not* verified is the unique identification

with 3D-vs-4D: the 4D theory has the same leading-order character expansion and would only diverge from 3D at higher orders or via the non-perturbative monopole density. In $D = 3$, the monopole gas is in a plasma phase at every β (Polyakov, BMK); in $D = 4$, it condenses at finite β_c . Distinguishing these requires either higher-order character corrections (technical, multi-week) or a direct monopole-density computation (well-defined but heavy, see §8).

The fourth witness for 3D-compatibility passes: **the strong-coupling character expansion gives an area law with $\sigma(\beta) > 0$ at every tested β , matching the Polyakov–BMK asymptotics for 3D compact $U(1)$ permanent confinement.** The verdict is structural compatibility, not unique identification with the 3D universality class.

6.3 NLO character expansion and structural cycle content

The leading-order area-law formula (19) is universal: any Wilson lattice gauge theory with the standard action gives $\langle W \rangle = (I_1/I_0)^A$ at leading order, regardless of the underlying graph topology and ambient dimension. The 3D-vs-4D distinction in compact $U(1)$ is fundamentally non-perturbative and enters at higher orders in the character expansion through corrections that depend on the plaquette-adjacency structure of the graph. A next-to-leading-order (NLO) follow-up sprint (XCWG-E, May 2026) computed these corrections on Rule B; this subsection reports its three substantive findings.

Setup. For Wilson action (6) and a Wilson loop $W(C) = e^{i\theta_C}$, the character expansion gives

$$\langle W(C) \rangle_\beta = \frac{N(w; \beta)}{Z(\beta)}, \quad N(w; \beta) = \sum_{n \in \mathbb{Z}^{|\mathcal{P}|}: d_1^\top n = w} \prod_P \frac{I_{n_P}(\beta)}{I_0(\beta)}, \quad (22)$$

with w the signed edge-indicator of C and the analogous $Z(\beta)$ summing over $d_1^\top n = 0$. LO is the minimum- $|n|_1$ assignment n_{LO} with $|n_{\text{LO}}|_1 = A_{\min}$, giving $(I_1/I_0)^{A_{\min}}$. NLO comes from $n = n_{\text{LO}} + z$ where z is a closed 2-cycle ($d_1^\top z = 0$), suppressed by an additional factor $(I_1/I_0)^{k_\delta}$ relative to LO, where k_δ is the smallest non-zero $|z|_1$ in the kernel of $d_1^\top$.

Smallest closed 2-cycle: $k_\delta = 3$ on Rule B. A BFS enumeration of connected plaquette subsets of size $k = 2, 3, \dots$ with sign patterns $z_P \in \{-1, 0, +1\}$ satisfying $d_1^\top z = 0$ returns $k_\delta = 3$ at both $n_{\max} = 2$ and $n_{\max} = 3$. A representative size-3 cycle at $n_{\max} = 2$ is $\{p_0, p_3, p_{38}\}$ with signs $(+1, -1, -1)$, where the three plaquettes pairwise share two edges and together cover six edges, each used by exactly two plaquettes. Geometrically this is a triangulated S^2 -like “triangle-prism” configuration: three L=4 plaquettes glued along common edge-axes, closing up into a 2-cycle.

For comparison, on the cubic lattice $\mathbb{Z}^4$ the smallest closed 2-surface is the boundary of a 3-cube made of 6 plaquettes ($k_\delta = 6$); on $\mathbb{Z}^3$ there are no closed 2-surfaces in the plaquette complex alone (3D confinement is monopole-driven rather than surface-driven). Rule B’s $k_\delta = 3$ is therefore structurally distinct from both: lower-order closed 2-surfaces than $\mathbb{Z}^4$, and the surfaces exist intrinsically in the L=4-plaquette 2-complex unlike $\mathbb{Z}^3$. The mechanism is direct: Rule B has unusually high plaquette density (~ 4.4 plaquettes per vertex at $n_{\max} = 2$, vs ~ 1 on $\mathbb{Z}^4$), so the kernel of $d_1^\top$ admits very low-order cycle relations. The second Betti number of the 2-complex is $b_2 = |\mathcal{P}| - \text{rank}(d_1) = 44 - 11 = 33$ at $n_{\max} = 2$ and $b_2 = 136$ at $n_{\max} = 3$ (with the 200-plaquette enumeration cap of the driver).

This is the principal new structural finding of XCWG-E: Rule B is a *non-cubic* graph in a strong combinatorial sense – its 2-complex carries closed 2-surfaces at lower order than any $\mathbb{Z}^d$ cubic lattice.

Direct NLO surface enumeration: truncation-sensitive. At $n_{\max} = 2$ with a single-plaquette Wilson loop ($A_{\min} = 1$) and NLO truncated to $|n|_1 \leq A_{\min} + k_\delta + 1 = 5$, the normalized $\langle W \rangle_{\text{NLO}}$ behaves as follows. The NLO string tension $\sigma_{\text{NLO}}(\beta) := -\ln \langle W \rangle_{\text{NLO}} / A$ stays below $\sigma_{\text{LO}}(\beta)$ at every tested β , consistent with NLO surfaces competing with the LO minimum and reducing the apparent string tension. Numerically σ_{NLO} crosses zero nominally around $\beta \sim 10$ –30, becoming slightly negative (-0.058 at $\beta = 100$).

This nominal zero crossing is accompanied by $\langle W \rangle_{\text{NLO}} > 1$ at large β (peaks at ~ 1.06), which is non-physical for a unitary Wilson loop. We diagnose this as a finite-cycle-enumeration *truncation artifact*, not a genuine 4D-style deconfinement transition. The mechanism: in the full (untruncated) sum, every higher-charge state with $I_n/I_0 \rightarrow 1$ at large β contributes proportionally to both N and Z , keeping their ratio bounded by 1. Truncating to $|n|_1 \leq \text{cutoff}$ breaks this symmetry asymmetrically because the count of valid surfaces with $d_1^\top n = w$ within the truncation budget grows differently from the count with $d_1^\top n = 0$. The pattern persists at $n_{\max} = 3$ with the same qualitative shape, confirming the artifact diagnosis.

Creutz-ratio cross-check: 3D-like at asymptotic area. The cleaner asymptotic-area diagnostic is the Creutz-like ratio of Wilson loops at adjacent areas:

$$\sigma_{\text{eff}}(A_1, A_2) := -\frac{\ln[\langle W(A_2) \rangle / \langle W(A_1) \rangle]}{A_2 - A_1}. \quad (23)$$

At LO this is independent of (A_1, A_2) and equals $\sigma_{\text{LO}}(\beta)$. At NLO, the small-area pair ($1 \rightarrow 2$) shows large NLO suppression, but the genuine asymptotic string tension is recovered by going to larger A . The $\sigma_{\text{eff}}(2 \rightarrow 3)$ pair at $n_{\max} = 2$ stays positive at every tested β and saturates at a positive plateau ~ 0.20 at large β , rather than crossing zero. The pair- $(3 \rightarrow 4)$ shows mild negative crossing at very large β , again consistent with the finite-cycle-budget truncation artifact above (small-cycle enumeration unable to fully match contributions to clusters of size 4 vs size 3).

The Creutz cross-check therefore reads 3D-like: *the asymptotic-area string tension stays positive at every β , with no clean deconfinement crossing*. This supports the leading-order area-law verdict (§6.2) at the next order accessible within the cycle budget, while honestly acknowledging that the single-plaquette diagnostic is dominated by truncation rather than physics.

Honest scope of the NLO sprint. The two NLO diagnostics *disagree* as numerical quantities: single-plaquette σ_{NLO} goes nominally negative at $\beta \sim 30$, while Creutz $\sigma_{\text{eff}}(2 \rightarrow 3)$ stays positive at all β . The reconciliation identified by Sprint XCWG-E is that the single-plaquette result is a truncation artifact (witnessed by $\langle W \rangle_{\text{NLO}} > 1$), while the Creutz analysis on asymptotic-area pairs is the structurally cleaner diagnostic, supporting the LO verdict. NLO at Rule B’s plaquette density (with $k_\delta = 3$) is too truncation-sensitive to give a clean 3D-vs-4D discriminator at the resolution this sprint accessed. *The four-witness leading-order verdict from §4–§6.2 stands unchanged; XCWG-E adds the $k_\delta = 3$ structural finding and a supporting (but not load-bearing) Creutz cross-check for the 3D-like reading at next order.* A clean NLO verdict would require either a much larger cycle-enumeration budget than this sprint’s, or a different observable that is less sensitive to surface-enumeration truncation – the monopole density (XCWG-F, §6.4).

6.4 Polyakov dilute monopole gas via DeGrand–Toussaint construction

The first four witnesses (§4–§6.2) each operate at leading order or at the linearised-Gaussian level. At those orders, the area-law formula $\langle W \rangle = (I_1(\beta)/I_0(\beta))^A$ is *universal* across compact $U(1)$ on

any graph carrying a 2-complex: it does not by itself distinguish 3D compact $U(1)$ (Polyakov 1977 permanent confinement) from 4D compact $U(1)$ (BMK 1977 / Guth 1980 deconfinement transition at $\beta_c \approx 1.01$). The 3D-vs-4D distinction is intrinsically *non-perturbative* and lives in the monopole sector: in 3D the magnetic monopoles form a dilute plasma with density exponentially suppressed but strictly nonzero at every $\beta > 0$, while in 4D the monopole density vanishes above β_c as the magnetic charges condense into a dilute gas admitting the deconfined Coulomb phase [13, 14, 15]. The canonical Polyakov-mechanism discriminator is therefore the existence and exponential decay of the monopole density $\rho_M(\beta)$, which the standard DeGrand–Toussaint construction [24] extracts from integer plaquette fluxes via discrete Stokes’ theorem on the dual 2-complex.

Why this matters beyond §6.1–§6.3. Witness 3 (§6.1) probes the linearised free-photon sector and loses the 2π periodicity that admits monopoles by construction; witness 4 (§6.2) is the universal area-law signature shared by 3D and 4D compact $U(1)$ at leading order; witness 5 (this subsection) probes the same compact integration domain that witness 4 ignores the periodicity of, and is the first diagnostic that can in principle distinguish 3D from 4D compact $U(1)$ rather than merely verify their shared leading-order behaviour. The XCWG-E NLO sprint (§6.3) attempted this distinction through higher-order character corrections but found single-plaquette σ_{NLO} truncation-dominated; the monopole density bypasses that obstruction.

DeGrand–Toussaint construction adapted to Rule B. The classical DeGrand–Toussaint procedure works on any oriented 2-complex. For a gauge configuration $\{\theta_e \in (-\pi, \pi]\}_{e \in E}$:

- (i) For each plaquette P , compute the oriented holonomy $\theta_P = \sum_{e \in \partial P} \text{sign}(P, e) \theta_e$. This is *not* reduced modulo 2π ; it can take any real value in $(-d_P\pi, d_P\pi]$ where $d_P = |\partial P| = 4$.
- (ii) Project to the principal branch $\bar{\theta}_P = \theta_P - 2\pi \text{nint}(\theta_P/2\pi) \in (-\pi, \pi]$; the integer *Dirac-string charge* on P is $m_P = (\theta_P - \bar{\theta}_P)/(2\pi) \in \mathbb{Z}$.
- (iii) For each elementary closed 2-surface S in the dual complex, the *monopole charge enclosed by* S is $m_S = \sum_{P \in S} \text{sign}(P, S) m_P \in \mathbb{Z}$, automatically integer by the discrete Stokes theorem on the integer-valued Dirac-string field. The quantity is gauge-invariant ($U(1)$ -modular in 2π) and depends only on the homology class of S .

On a cubic lattice, the elementary closed 2-surfaces are 3-cube boundaries (6 plaquettes each). On Rule B, the analog is supplied directly by the XCWG-E sprint (§6.3): the smallest closed 2-cycle has size $k_\delta = 3$, given by triangle-prism configurations of three $L = 4$ plaquettes pairwise sharing two edges. Enumerating all $\{-1, 0, +1\}$ -valued closed 2-cycles of size 3 returns 60 elementary monopole sites at $n_{\text{max}} = 2$ and 80 at $n_{\text{max}} = 3$. Higher-size cycles (sizes 4–6) also exist (557 sites of size ≤ 6 at $n_{\text{max}} = 2$); these are higher order in the dilute-gas expansion and are not used as elementary sites here.

The monopole density observable is

$$\rho_M(\beta) = \frac{1}{|\mathcal{S}_{\text{elem}}|} \left\langle \sum_{S \in \mathcal{S}_{\text{elem}}} |m_S| \right\rangle_\beta, \quad (24)$$

where the expectation is taken with respect to the Boltzmann measure for the compact $U(1)$ Wilson action (6) at coupling β .

Monte Carlo of compact $U(1)$ Wilson on Rule B. We use standard single-link Metropolis–Hastings on the Wilson action $S_W = \beta \sum_P (1 - \cos \theta_P)$: for each edge e , propose $\theta_e \rightarrow \theta_e + \delta$ with $\delta \sim \text{Unif}(-\delta_{\max}, \delta_{\max})$ and accept with probability $\min(1, e^{-\Delta S_W})$, then fold back to $(-\pi, \pi]$. The proposal width $\delta_{\max}$ is auto-tuned to acceptance rate ~ 0.5 (approaching π at small β , dropping to ~ 0.18 at large β). At $n_{\max} = 2$ we run $n_{\text{therm}} = 2000$ thermalisation sweeps and $n_{\text{sample}} = 1000$ samples with $\Delta_{\text{sweep}} = 40$ between samples (40,000 sampling sweeps per β), at 15 values of $\beta \in [0.05, 30]$; at $n_{\max} = 3$, $n_{\text{therm}} = 800$, $n_{\text{sample}} = 400$, $\Delta_{\text{sweep}} = 30$ at 10 values of β . Block standard errors are computed from 10 contiguous sample blocks. Total wall time ~ 3 minutes.

The integer-valued character of m_S implies a Monte Carlo detection floor: $\rho_M^{\text{floor}} = 1/(n_{\text{sample}}|\mathcal{S}_{\text{elem}}|)$, equal to 1.67×10^{-5} at $n_{\max} = 2$ and 3.13×10^{-5} at $n_{\max} = 3$. A measured $\rho_M = 0$ at large β means “below MC detection,” not “identically zero.”

Numerical results. At $n_{\max} = 2$:

β	$\rho_M(\beta)$	$\pm \sigma$ (block)
0.05	2.373×10^{-1}	2.18×10^{-3}
0.10	2.133×10^{-1}	2.61×10^{-3}
0.20	1.592×10^{-1}	2.60×10^{-3}
0.30	8.035×10^{-2}	3.73×10^{-3}
0.50	7.40×10^{-3}	8.39×10^{-4}
0.70	8.00×10^{-4}	2.69×10^{-4}
1.00	5.00×10^{-5}	5.00×10^{-5}
≥ 1.5	0 (below floor 1.67×10^{-5})	—

At $n_{\max} = 3$:

β	$\rho_M(\beta)$	$\pm \sigma$ (block)
0.05	2.362×10^{-1}	3.47×10^{-3}
0.10	2.127×10^{-1}	5.08×10^{-3}
0.20	1.519×10^{-1}	4.91×10^{-3}
0.30	5.531×10^{-2}	2.55×10^{-3}
0.50	3.50×10^{-3}	4.72×10^{-4}
0.70	6.875×10^{-4}	1.02×10^{-4}
1.00	9.375×10^{-5}	4.77×10^{-5}
≥ 2.0	0 (below floor 3.13×10^{-5})	—

The dynamic range of nonzero ρ_M spans *five orders of magnitude* at both cutoffs (2.4×10^{-1} at $\beta = 0.05$ to $\sim 5 \times 10^{-5}$ at $\beta = 1.0$), with monotonic decrease throughout. The two cutoffs agree at small β to within MC error (e.g. $\rho_M(0.1) \approx 0.21$ at both), reflecting that triangle-prism monopole sites are *local* objects whose per-site density does not scale anomalously with graph size.

Polyakov dilute-gas fit. Polyakov 1977 predicts the dilute monopole-plasma form $\rho_M(\beta) \sim A e^{-c\beta}$, with $c = 2\pi^2 v_0$ where v_0 is the cell-volume scale of the dual lattice. Weighted least-squares fit of $\log \rho_M(\beta) = \log A - c\beta$ on the range where $\rho_M > 0$ ($\beta \in [0.05, 1.0]$):

$n_{\max}$	A	c	R^2	$\beta_{\text{floor}}^{\text{pred}}$
2	0.747	9.40	0.981	1.14
3	0.528	8.95	0.981	1.09

Two structural facts hold uniformly. *First*, the exponent c is $n_{\max}$ -independent within 5% ($c_2/c_3 = 1.05$, well within MC noise), which is a genuine 3D-compatible signature: in the dilute regime, local physics observables should not scale anomalously with V or E , and the exponential decay rate is a property of the dual cell-volume, not the finite-graph cutoff. *Second*, the fit-predicted floor crossover $\beta_{\text{floor}}^{\text{pred}} = \log(A/\rho_M^{\text{floor}})/c$ lands at $\beta \approx 1.1$ at both cutoffs, which is consistent with the observed $\beta_{\text{zero-onset}} = 1.0\text{--}1.5$ within the resolution of the β grid. The apparent “zeros” at $\beta \geq 1.5$ are therefore *exponentially small positive values below MC detection*, not genuine zero crossings: the Polyakov dilute-plasma signature, not a 4D-style condensation transition.

Identifying $c = 2\pi^2 v_0$ on Rule B gives $v_0^{\text{Rule B}} = c/(2\pi^2) \approx 9.20/19.74 \approx 0.466$, of order unity as expected. The numerical v_0 differs from the cubic-lattice baseline because Rule B is irregular and non-cubic; the structural-compatibility statement is that the *form* of the decay (exponential with $\mathcal{O}(1)$ exponent) is correct, not that the numerical v_0 should match the cubic case. A sharper match would require analytical computation of the dual cell-volume on Rule B and is deferred (§8).

Pass criteria for the fifth witness. The Polyakov dilute-plasma diagnostic passes if:

- (a) $\rho_M(\beta) > 0$ wherever it can be resolved above the MC detection floor: **PASSES** (five decades of nonzero ρ_M across $\beta \in [0.05, 1.0]$ at both $n_{\max} \in \{2, 3\}$);
- (b) the nonzero values are monotonically decreasing in β : **PASSES**;
- (c) the decay is exponential with $c > 0$ and $R^2 \gtrsim 0.95$: **PASSES** ($c \approx 9.2$ across both cutoffs, $R^2 = 0.981$);
- (d) the apparent “zeros” at large β are consistent with ρ_M having dropped below the MC detection floor rather than constituting a genuine zero crossing: **PASSES** ($\beta_{\text{floor}}^{\text{pred}} \approx 1.1$ matches observed $\beta_{\text{zero-onset}} = 1.0\text{--}1.5$);
- (e) the decay rate c is consistent across $n_{\max}$ (no anomalous size scaling): **PASSES** ($c_2/c_3 = 1.05$).

Honest scope. The MC sample sizes give a detection floor $\rho_M \gtrsim 1.7 \times 10^{-5}$. Predicted ρ_M at $\beta = 2$ from the fit is $\sim 0.747 e^{-18.8} \sim 5 \times 10^{-9}$, four orders below floor; resolving nonzero ρ_M at $\beta = 2$ would require $\gtrsim 10^4$ longer MC runs. The conclusion that $\rho_M(\beta) > 0$ at $\beta \in [1.5, \infty)$ therefore rests on exponential extrapolation of a fit that holds to $R^2 = 0.981$ over five decades. We do not directly measure the thermodynamic limit; what is probed is consistency between $n_{\max} = 2$ and $n_{\max} = 3$ (c stable within 5%). The Wilson action $\beta(1 - \cos \theta_P)$ is used rather than the Manton or Villain variants; this is the conventional choice at the discriminator level and matches Polyakov 1977.

Verdict. The fifth witness passes: the monopole density on Rule B’s compact $U(1)$ Wilson theory exhibits the Polyakov 3D dilute-plasma signature rather than the 4D condensation transition, with five decades of exponential decay $\rho_M \sim A e^{-c\beta}$ at $R^2 = 0.981$, $n_{\max}$ -independent decay rate $c \approx 9.2$ within 5%, and no zero crossing visible above the MC detection floor. The fifth witness is structurally distinct from the first four: it is the only one that probes the non-perturbative monopole sector directly, and is the canonical Polyakov discriminator between 3D and 4D compact $U(1)$.

6.5 Full Monte Carlo Wilson loops

Witnesses 3 and 4 (§6.1 and §6.2) operate at leading order in two different sectors: the linearised Gaussian free-photon sector and the strong-coupling small- β character expansion. Witness 5 (§6.4) probes the non-perturbative monopole sector via Metropolis–Hastings MC of the same compact Wilson action that underlies witness 4. A natural sixth probe is to evaluate Wilson loops themselves on the full compact $U(1)$ measure – directly testing whether the area-law mechanism observed at LO (witness 4) survives in the non-perturbative regime, and whether the Polyakov dilute-gas prediction $\sigma(\beta) \sim \exp(-(c_\rho/2)\beta)$ relating witness 4’s string tension to witness 5’s monopole density holds quantitatively on the irregular non-cubic Rule B graph. Sprint XCWG-G (May 2026) executed this probe at $n_{\max} \in \{2, 3\}$ using the same Metropolis–Hastings infrastructure as XCWG-F (§6.4).

Methodology. The compact $U(1)$ Wilson measure $P(\{\theta_e\}) \propto \exp(-S_W)$ with action (6) is sampled via single-link Metropolis updates of link angles $\theta_e \in (-\pi, \pi]$. Proposal width $\delta_{\max}$ is auto-tuned to acceptance $\sim 50\%$; thermalisation is 2000 sweeps with samples taken every 20 sweeps. We use $n_{\text{sample}} = 2000$ at $n_{\max} = 2$ and $n_{\text{sample}} = 600$ at $n_{\max} = 3$. For each sample $\{\theta_e\}$ and each area-controlled Wilson loop C with signed-edge vector σ_C , we compute $W(C; \theta) = \cos(\sigma_C \cdot \theta)$ and average over MC samples.

Distinct from witness 4 (XCWG-D), which used the leading-order character expansion $\langle W(C) \rangle_\beta = (I_1(\beta)/I_0(\beta))^{A(C)}$, witness 6 (XCWG-G) samples the full compact measure at every β , including all higher orders of the character expansion AND all monopole fluctuations simultaneously. Distinct from witness 3 (XCWG-C), which used the linearised quadratic form $S(W) = \langle W, K^+ W \rangle$ blind to the 2π periodicity, witness 6 sees the monopole-mediated disorder-disorder correlations that drive the 3D compact $U(1)$ confinement in continuum.

Area-controlled loop construction. We reuse XCWG-D’s `grow_area_A_clusters` to build clusters of A plaquettes whose oriented boundary is a single closed loop. At $n_{\max} = 2$, available perimeters by area are $A = 1 : L \in \{4\}$, $A = 2 : L \in \{4, 6\}$, $A = 3 : L \in \{4, 6, 8, 10\}$, $A = 4 : L \in \{4, 6, 8, 10, 12\}$, $A = 5 : L \in \{4, 6, 8, 10, 12\}$. We select 10 loops per area, round-robin across perimeters. At $n_{\max} = 3$, $L = 4$ extends only to $A \in \{1, 2, 3\}$, and $L \in \{10, 12\}$ no longer appear – a finite-volume constraint that limits the fixed- L tests at the larger cutoff.

Numerical results: $\sigma_{\text{ens}}(\beta)$ at $n_{\max} = 2$. A naive ensemble-mean area-law fit $\log \langle W(A) \rangle = -\sigma A + c$ on per-area ensemble means gives $\sigma_{\text{ens}}(\beta) > 0$ at every $\beta \in [0.1, 10]$, monotone decreasing in β :

β	σ_{ens}	σ_{comb}	μ_{comb}	σ_{LO} (XCWG-D)
0.1	0.377 ± 0.074	-0.007 ± 0.016	+0.342	2.997
0.3	0.174 ± 0.009	-0.006 ± 0.003	+0.217	1.908
1.0	0.033 ± 0.001	-0.002 ± 0.000	+0.034	0.807
3.0	0.011 ± 0.000	-0.000 ± 0.000	+0.010	0.211
10.0	0.003 ± 0.000	-0.000 ± 0.000	+0.003	0.053

The naive σ_{ens} is positive, monotone decreasing, and has the qualitative shape that a 3D compact $U(1)$ confining theory would produce. The ratio $\sigma_{\text{ens}}/\sigma_{\text{LO}}$ at small β is 0.05–0.13, well below 1 – the leading-order XCWG-D prediction overshoots the full-MC ensemble fit by an order of magnitude, signalling that the LO formula overcounts the area law on this graph.

Area-vs-perimeter separation: the load-bearing finite-volume finding. The leading-order character expansion has no perimeter dependence at fixed area: $\langle W \rangle_{\text{LO}} = (I_1/I_0)^A$ depends only on A , not on L . At full MC, this prediction is *falsified* for Rule B at the available graph sizes. Per-loop data within an area class shows strong perimeter dependence: at $\beta = 1$, area $A = 4$ (five distinct perimeters):

L	$\langle W \rangle$ at $\beta = 1, A = 4$
4	0.878 ± 0.002
6	0.826 ± 0.002
8	0.774 ± 0.002
10	0.706 ± 0.002
12	0.687 ± 0.002

The spread across perimeters at fixed A is 0.19 units of W , which is $20\text{--}100\times$ larger than the MC standard error on the per-area ensemble mean. A joint weighted least-squares fit $\log W_i = -\sigma A_i - \mu L_i + c$ on per-loop data disentangles area and perimeter contributions: the genuine area-law coefficient σ_{comb} is *statistically zero* at every β , and the apparent confinement signal is carried by the perimeter coefficient $\mu_{\text{comb}} \approx \sigma_{\text{ens}}$. A fixed-perimeter restriction $\sigma_{L=4}$ at $\beta = 1$ gives -0.003 , confirming that the per-area slope vanishes when L is held fixed.

The mechanism is geometric: in the area-controlled cluster enumeration, larger A comes with longer accessible L (a roughly linear relation $\partial L/\partial A \sim 1$ at the cluster sizes available on Rule B at $n_{\text{max}} \leq 3$), so the per-loop $\log \langle W \rangle \propto -\mu L$ appears in the per-area fit as $-\sigma_{\text{ens}} A$. **On Rule B at $n_{\text{max}} \in \{2, 3\}$, full-MC Wilson loops are perimeter-dominated, not area-dominated.** This is a genuine finite-volume effect that scales with V_B : as the graph grows, the available cluster perimeters at fixed area should saturate (the graph diameter caps L at fixed A), allowing a clean separation in the joint fit.

Polyakov dilute-gas cross-check: exponential shape, wrong rate. Polyakov 1977 predicts the dilute-monopole-gas relation $\sigma_{\text{Polyakov}}(\beta) \propto \sqrt{\rho_M(\beta)/\beta} \propto \exp(-(c_\rho/2)\beta)$, so the predicted exponential rate constant in $\log \sigma$ vs β is $c_\sigma^{\text{predicted}} = c_\rho/2 = 4.70$ (using XCWG-F’s $c_\rho = 9.40$ at $n_{\text{max}} = 2$). Fitting $\log \sigma_{\text{ens}}(\beta)$ against β over $\beta \in [0.1, 3.0]$:

Cutoff	c_σ^{MC}	$c_\sigma/(c_\rho/2)$	R^2 of exp fit
$n_{\text{max}} = 2$	1.20	0.26	0.83
$n_{\text{max}} = 3$	1.63	0.35	0.88

The measured rate constant is a *factor of 3–4× smaller than predicted* at both cutoffs. The exponential shape itself fits with $R^2 = 0.83\text{--}0.88$ (not great, not terrible: the data is exponentially decaying at all tested β), but the quantitative agreement is poor. Pushing $n_{\text{max}} = 2 \rightarrow 3$ improves the agreement modestly ($0.26 \rightarrow 0.35$); extrapolating linearly suggests $n_{\text{max}} \sim 7$ might recover the textbook rate, but the extrapolation is not robust.

Reading: non-cubic dual-lattice topology, not a contradiction of XCWG-F. The Polyakov derivation assumes (a) the 3D-cubic dual-lattice structure “dual lattice = points with monopoles on them”; (b) the integer-valued Dirac-string field m_P has a continuum-like Coulomb-gas free energy that scales as ρ_M ; (c) the area law inherits the $\sqrt{\rho_M/\beta}$ scaling from the correlation function of two

test charges in the monopole plasma. On Rule B at $n_{\max} = 2$, the dual lattice is far from regular cubic: the closed 2-surfaces on which monopole charges sit are size-3 triangle-prism configurations (§6.3, $k_\delta = 3$ vs $\mathbb{Z}^4$'s $k_\delta = 6$), and there are 60 such elementary monopole sites packed onto 44 plaquettes – a mean of ~ 4.4 plaquettes per vertex, far above the ~ 1 of $\mathbb{Z}^4$. Steps (b) and (c) of the Polyakov derivation likely fail in form on this graph: the standard sine-Gordon dualization is specific to $\mathbb{Z}^d$ topology, and the “test-charge correlator” inherits the perimeter-law contamination identified above.

The XCWG-F monopole-density result $\rho_M(\beta) > 0$ at all tested β stands as a measured fact (the integer flux extraction is a geometric quantity that does not depend on the dilute-gas assumptions). What XCWG-G reveals is that the *Polyakov-style derivation of σ from ρ_M via the dilute-gas formula does not transfer to the Rule B finite graph in the form the textbook 3D continuum result would predict*. The exponential shape of $\sigma_{\text{ens}}(\beta)$ is qualitatively the right Polyakov signature; the rate constant disagreement is the natural disclaimer on a small discrete graph with mixed-density adjacency.

Verdict: weak pass with finite-volume scope. The XCWG-G test consists of three sub-claims: (1) $\sigma_{\text{MC}}(\beta) > 0$ at all tested β : *YES* for the naive ensemble fit σ_{ens} , *NO* for the joint σ_{comb} and fixed- L $\sigma_{L=4}$. (2) $\sigma_{\text{MC}}(\beta)$ monotonically decreasing in β : *YES* for σ_{ens} . (3) Either $\sigma_{\text{MC}} \approx \sigma_{\text{LO}}$ at small β (XCWG-D consistency) or $c_\sigma^{\text{MC}} \approx c_\rho/2$ (Polyakov / XCWG-F consistency): *NEITHER* holds quantitatively ($\sigma_{\text{ens}}/\sigma_{\text{LO}} \sim 0.05\text{--}0.13$, $c_\sigma/(c_\rho/2) \sim 0.26\text{--}0.35$).

A *clean pass* would require both (1) joint-fit area-law positivity *and* (2) quantitative Polyakov-rate agreement. Neither holds at the graph sizes accessible in this sprint. A *weak pass* in the perimeter-dominated finite-volume regime holds: σ_{ens} is monotone-decreasing positive at every β , matching the structural expectation of 3D compact $U(1)$ qualitatively (and the previous witnesses XCWG-D and XCWG-F directionally), but the data is perimeter-dominated and a clean separation of area-law from perimeter-law content would require larger Rule B graphs.

The sixth witness passes weakly with explicit finite-volume scope: the full-MC Wilson loops on Rule B at $n_{\max} \leq 3$ are perimeter-dominated, with $\sigma_{\text{ens}}(\beta) > 0$ monotone-decreasing qualitatively consistent with 3D compact $U(1)$. The Polyakov dilute-gas exponential shape is recovered; the rate constant is 3–4 $\times$ off from XCWG-F’s prediction, an observation that v5 reframes (see §6.7 below). The sixth witness is structurally compatible with the 3D-compact- $U(1)$ reading at the level it accesses, with an honest finite-volume caveat that the available graph sizes do not yet permit a clean area-law extraction.

$n_{\max} = 4$ extension (XCWG-G, v5 update). The v4 expectation was that area-vs-perimeter separation would clean up at the larger graph ($V = 60$, $E = 312$, $\beta_1 = 253$). The actual finding is more interesting and more structural: $\sigma_{\text{ens}}(\beta)$ remains positive and monotonically decreasing (+0.172 at $\beta = 0.3$, +0.040 at $\beta = 1.0$, +0.011 at $\beta = 3.0$), confirming the qualitative confinement signal at the larger cutoff. But the joint $\sigma A + \mu L$ fit gives σ_{comb} that moves *deeper negative* with significance, not weaker: $\sigma_{\text{comb}} = -0.080 \pm 0.002$ ($z = -48$) at $\beta = 0.3$, -0.021 ± 0.001 ($z = -26$) at $\beta = 1.0$, and -0.008 ± 0.000 ($z = -39$) at $\beta = 3.0$. The trend across cutoffs $\sigma_{\text{comb}}/\text{SE}$ goes $-2.0 \rightarrow -4.8 \rightarrow -48$ from $n_{\max} = 2 \rightarrow 3 \rightarrow 4$ at $\beta = 0.3$, in the wrong direction for finite-volume scaling.

The mechanism is identified in the cluster construction itself: at $n_{\max} = 4$ the available-perimeter table for area-controlled clusters is $L_{\min}(A) = 4, 4, 6, 8, 12, 10$ for $A = 1, 2, 3, 4, 5, 6$, which is *non-monotonic* – jumping from $L_{\min}(A = 4) = 8$ to $L_{\min}(A = 5) = 12$ then back to $L_{\min}(A = 6) = 10$. At $\beta = 0.3$ this produces $\langle W(A = 4) \rangle = 0.578 > \langle W(A = 3) \rangle = 0.515$ (a

20σ effect) and $\langle W(A=6) \rangle = 0.406 > \langle W(A=5) \rangle = 0.354$ – anti-area-law in the ensemble, but driven by the cluster’s $A \leftrightarrow L_{\min}$ coupling, not by gauge dynamics. The clean fixed- $L=6$ slice at $\beta=0.3$ gives $\sigma_{\text{fix-}L=6} = -0.004 \pm 0.003$ (consistent with zero, $z = -1.4$), which is the honest area-law coefficient at the available MC statistics. This is fully consistent with the Polyakov-dilute-gas prediction $\sigma_{\text{Polyakov}} \sim 5 \times 10^{-9}$ at the same β , four orders of magnitude below the MC detection floor. The diagnosis therefore sharpens from “finite-volume artifact, fades at larger $n_{\max}$ ” to “cluster-topology-induced non-linear $A \leftrightarrow L_{\min}$ coupling that the joint linear fit cannot resolve at any tested graph size; the honest area-law signal is at or below MC noise.” The next sharpening requires either (a) extended MC statistics to push the detection floor below the Polyakov-predicted σ , or (b) a fixed- L ensemble construction by design rather than cluster-area control, which would isolate the area-coefficient from the cluster geometry.

6.6 Polyakov loops via Euclidean-time compactification

Witnesses 1–6 are all zero-temperature observables. A natural seventh diagnostic probes finite-temperature behaviour: does the Rule B Wilson construction exhibit a Kaluza-Klein deconfinement transition as the temperature is raised, or does it stay confined at all T as 3D compact $U(1)$ does? This distinction is the canonical 3D-vs-4D finite-temperature discriminator (Svetitsky-Yaffe universality in 4D, no transition in 3D).

The Polyakov loop – the canonical deconfinement order parameter at finite T – requires a distinguished temporal direction. Rule B as a static spatial graph does not provide one; the natural construction is the *product graph* $G_B \times C_{N_t}$, where C_{N_t} is the cyclic chain of length N_t playing the role of the Euclidean-time circle at inverse temperature $\beta_{\text{thermal}} \propto N_t$. This is the discrete analog of taking $\mathbb{R}^3 \times S^1_\beta$ in the continuum, specialised to the irregular Rule B spatial graph rather than the round S^3 . Sprint XCWG-H (May 2026) built this construction from scratch and executed first-pass Metropolis MC on the product graph.

Construction of $G_B \times C_{N_t}$. The product graph has vertices (v, t) with $v \in V_B$ and $t \in \{0, 1, \dots, N_t - 1\}$, and two edge classes: *spatial* edges $((v, t), (v', t))$ for each $\{v, v'\} \in E_B$ at each t , and *temporal* edges $((v, t), (v, t+1 \bmod N_t))$ for each $v \in V_B$ at each t . Plaquettes are *spatial* (one per Rule B plaquette at each t) or *temporal-rectangular* (one per spatial edge between consecutive t -slices). Bookkeeping at $n_{\max} = 2$ (with $V_B = 10$, $E_B = 20$, $P_B = 44$):

Quantity	Formula	$N_t = 2$	$N_t = 4$	$N_t = 8$
V	$V_B \cdot N_t$	20	40	80
E_{spatial}	$E_B \cdot N_t$	40	80	160
E_{temporal}	$V_B \cdot N_t$ (V_B at $N_t = 2$)	10	40	80
E total	sum	50	120	240
P_{spatial}	$P_B \cdot N_t$	88	176	352
P_{temporal}	$E_B \cdot N_t$ (E_B at $N_t = 2$)	20	80	160
P total	sum	108	256	512
$\beta_1 = E - V + 1$		31	81	161
$\dim \ker(d_1)$		0	1	1

The $N_t = 2$ case is degenerate: C_2 has 2 vertices but only 1 undirected edge per spatial site, so $E_{\text{temporal}} = V_B$ rather than $V_B \cdot N_t$, and similarly each rectangular plaquette at $N_t = 2$ collapses to a single plaquette traversed forward or backward, so $P_{\text{temporal}} = E_B$ rather than $E_B \cdot N_t$. This degeneracy has a direct consequence for the Polyakov-loop observable, identified below.

Homological Polyakov class: $\dim \ker(d_1) = 1$ **as a structural prediction.** The cocycle-defect column $\dim \ker(d_1) = E - V + 1 - \text{rank}(d_1)$ is the dimension of the first cohomology that is *not* generated by plaquette boundaries. For $N_t \in \{4, 8\}$ this is exactly **1**, and the one extra cocycle is precisely the *Polyakov-loop class* – the gauge-invariant winding around the temporal cycle C_{N_t} . This is the homological signature that makes the Polyakov loop a well-defined observable on the product graph: it lives in a class that plaquettes alone do not cover, so the temporal-winding holonomy survives Gauss-law gauge equivalence. **The Polyakov loop is the right observable to probe finite- T deconfinement on $G_B \times C_{N_t}$ not by analogy to continuum 4D U(1), but as a homologically forced observable of the discrete object.**

At $N_t = 2$, $\dim \ker(d_1) = 0$: the would-be Polyakov class is generated by plaquette boundaries and is therefore not an independent observable. This is consistent with the kinematic identity for the Polyakov loop at $N_t = 2$, below.

Wilson action and Polyakov loop definition. The Wilson U(1) action on the product graph sums over both spatial and temporal-rectangular plaquettes,

$$S_W[\theta] = \beta \sum_{P \in \mathcal{P}} (1 - \cos \theta_P), \quad \theta_P = \sum_{e \in \partial P} \text{sign}(P, e) \theta_e, \quad (25)$$

with $\mathcal{P} = \mathcal{P}_{\text{spatial}} \cup \mathcal{P}_{\text{temporal}}$. We take $\beta_{\text{spatial}} = \beta_{\text{temporal}} = \beta$ for the first pass; an anisotropic version with $\beta_t \neq \beta_s$ is a natural future direction (Hamiltonian-limit formulation of finite- T lattice gauge theory).

The Polyakov loop at spatial vertex $v \in V_B$ is the path-ordered product of *temporal* link variables around the cycle C_{N_t} anchored at v ,

$$P(v) = \prod_{t=0}^{N_t-1} U_{(v,t) \rightarrow (v,t+1 \bmod N_t)} = \exp\left(i \sum_{t=0}^{N_t-1} \text{sign}_{v,t} \theta_{e_{v,t}}\right), \quad (26)$$

where $e_{v,t}$ is the (canonically oriented) temporal edge between (v, t) and $(v, t + 1 \bmod N_t)$ and $\text{sign}_{v,t} \in \{\pm 1\}$ is the orientation sign. The observable $\langle P \rangle = \langle \frac{1}{V_B} \sum_v P(v) \rangle$ is the gauge-invariant Polyakov-loop expectation; $\langle P \rangle = 0$ means confined, $\langle P \rangle \neq 0$ means deconfined.

The $N_t = 2$ kinematic degeneracy. At $N_t = 2$ the temporal cycle collapses to a single edge per spatial vertex. The Polyakov path at v is $P(v) = U_{(v,0) \rightarrow (v,1)} \cdot U_{(v,1) \rightarrow (v,0)} = e^{i\theta_e} \cdot e^{-i\theta_e} = 1$ *identically*, for every gauge configuration. This is a kinematic identity, not a deconfinement signal: the “loop” backtracks on a single edge and carries no net winding. The MC output at $N_t = 2$ is therefore $\langle P \rangle = 1.0000$ exactly (no noise) at every β . We exclude $N_t = 2$ from the confinement verdict; the first nontrivial N_t is 3, and $N_t \in \{4, 8\}$ are the operating points. This degeneracy is detected by $\dim \ker(d_1) = 0$ at $N_t = 2$.

Monte Carlo results. Single-link Metropolis on action (25) with auto-tuned δ_{max} targeting $\sim 50\%$ acceptance; 2000 thermalisation sweeps from cold start $\theta = 0$, then 800 samples taken every 20 sweeps (16000 sampling sweeps per cell). Block-bootstrap errors with 10 blocks of 80 samples. Panel: $(\beta, N_t) \in \{0.3, 1.0, 3.0\} \times \{4, 8\}$.

β	N_t	$\text{Re} \langle P \rangle$	$\text{Im} \langle P \rangle$	$ \langle P \rangle $
0.3	4	-0.003 ± 0.006	-0.004 ± 0.009	0.005 ± 0.004
1.0	4	$+0.012 \pm 0.017$	$+0.002 \pm 0.012$	0.012 ± 0.009
3.0	4	-0.062 ± 0.053	$+0.036 \pm 0.047$	0.072 ± 0.043
0.3	8	-0.008 ± 0.011	$+0.003 \pm 0.007$	0.009 ± 0.006
1.0	8	$+0.008 \pm 0.009$	-0.001 ± 0.012	0.008 ± 0.010
3.0	8	-0.023 ± 0.024	$+0.029 \pm 0.038$	0.037 ± 0.020

All six $|\langle P \rangle|$ values are within 2σ of zero, the largest being 0.072 ± 0.043 at $(\beta, N_t) = (3, 4)$ – a 1.7σ excursion, well inside MC noise for 800 samples. Verdict at each cell: *CONFINED*. The pattern is robust between $N_t = 4$ and $N_t = 8$ (the two N_t values bracket the temperature range while keeping the spatial construction fixed).

Two-point Polyakov correlator: exponential decay with the correct N_t -dependence. Binning the correlator $|\langle P(v)P^*(w) \rangle|$ by spatial graph distance $d(v, w) \in \{1, 2, 3, 4\}$ on G_B (diameter 4 at $n_{\max} = 2$), at $\beta = 3$:

N_t	$d = 1$	$d = 2$	$d = 3$	$d = 4$	per-step ratio
4	9.5×10^{-4}	8.1×10^{-4}	6.7×10^{-4}	4.9×10^{-4}	~ 0.83
8	6.2×10^{-4}	4.7×10^{-4}	3.4×10^{-4}	2.2×10^{-4}	~ 0.72

Two clean structural facts. *First*, the correlator decays approximately exponentially with spatial graph distance at fixed N_t , consistent with confined-phase quark-antiquark string tension. *Second*, the per-step decay ratio at $N_t = 4$ is *larger* than at $N_t = 8$, which is the right direction for hotter (smaller N_t) systems having longer spatial correlation lengths: the quark-antiquark mass $m_{QQ} \sim \sigma \cdot N_t$ shrinks with N_t . The dynamic range across 4 graph steps is modest (~ 5 – $8\times$), capped by the spatial graph diameter being only 4 at $n_{\max} = 2$; a clean string-tension extraction would require $n_{\max} = 3$ (diameter 6) or larger and is deferred (§8).

Comparison to textbook 3D compact $U(1)$ at finite T . The expectation $\langle P \rangle = 0$ persistent at all T in 3D compact $U(1)$ is a corollary of two classical results. (a) *Polyakov 1977* [13]: the 3D compact $U(1)$ theory is permanently confined at all β via monopole condensation; the dual sine-Gordon theory has a mass gap for any $\beta > 0$. At finite T , the Kaluza-Klein reduction $S_\beta^1 \rightarrow$ effective 2D theory retains a mass gap. (b) *Wilson 1974* [12] §5: 2D $U(1)$ lattice gauge is confined trivially (the partition function factorises per plaquette, the Wilson loop has exact area law at all β via single-plaquette character integration). The conjunction of (a) and (b) means: as we raise T on the 3D theory (shrink N_t), we KK-reduce toward 2D $U(1)$, which is *more* confined, not less. **No deconfinement transition is possible.**

This is the structural prediction we test. Our $N_t \in \{4, 8\}$ data at $\beta \in \{0.3, 1.0, 3.0\}$ measures $|\langle P \rangle|$ within 2σ of zero at every cell. *Consistent with the 3D-compact- $U(1)$ persistent- confinement prediction.* For contrast: 4D compact $U(1)$ has the BMK 1977 [14] / Guth 1980 [15] transition at $\beta_c \approx 1.01$ at $T = 0$, and a Svetitsky-Yaffe deconfinement transition at finite T above some N_t -dependent threshold. Our Rule B data does *not* see this transition, consistent with the prior six witnesses pointing to 3D rather than 4D.

Verdict: clean pass. The XCWG-H test consists of four sub-claims: (1) the product graph $G_B \times C_{N_t}$ admits a well-defined Polyakov-loop observable: *YES* (homological Polyakov class $\dim \ker(d_1) = 1$ at $N_t \in \{4, 8\}$). (2) The kinematic degeneracy at $N_t = 2$ is correctly identified and excluded: *YES* ($\dim \ker(d_1) = 0$ matches $P \equiv 1$ identity). (3) $\langle P \rangle = 0$ within 3σ at every non-degenerate (β, N_t) : *YES* (all 6 cells within 2σ , largest 1.7σ). (4) Two-point correlator exhibits confined-phase exponential decay with N_t -dependence in the right direction: *YES* (per-step ratio 0.83 at $N_t = 4$ vs 0.72 at $N_t = 8$, both at $\beta = 3$).

The seventh witness passes cleanly: the Polyakov loop on $G_B \times C_{N_t}$ at $N_t \in \{4, 8\}$ is consistent with zero at all tested β , with the two-point correlator decaying exponentially with spatial graph distance and the decay rate increasing with N_t , exactly the persistent-confinement signature of 3D compact $U(1)$ at finite temperature. The seventh witness adds the structurally distinct finite- T probe to the six zero-temperature witnesses, and the verdict is consistent: Rule B’s Wilson construction does *not* exhibit the Kaluza-Klein deconfinement transition that 4D compact $U(1)$ does, and the homological mechanism (Polyakov class as the one extra first-cohomology cocycle) is the structurally novel feature that makes the observable well-defined on the product graph in the first place.

6.7 Polyakov-rate resolution and non-cubic dual-lattice structure

The v4 paper recorded an honest-scope caveat (vi): the Polyakov dilute-gas relation $\sigma(\beta) = (4/\pi)\sqrt{\rho_M/\beta_V}$ predicts rate constant $c_\sigma = c_\rho/2 = 4.70$ from XCWG-F’s $c_\rho = 9.40$, while XCWG-G measured $c_\sigma^{\text{MC}} = 1.20$ at $n_{\text{max}} = 2$ – a factor-of-3.9 disagreement. The v4 framing attributed this to non-cubic dual-lattice corrections, but deferred the structural-correction calculation. Sprint XCWG-J (May 2026) builds the dual graph and resolves the disagreement: it is not a missing rate correction, but a category mismatch between two different observables.

The dual graph G_B^* at $n_{\text{max}} = 2$. Vertices of G_B^* are Rule B’s plaquettes (44 at $n_{\text{max}} = 2$); edges connect plaquettes sharing an edge in Rule B. The construction gives $V_{\text{dual}} = 60$, $E_{\text{dual}} = 288$, average degree 9.60 (vs $\mathbb{Z}^3$ ’s constant degree 6). The dual Laplacian $L_{\text{dual}} = D_{\text{dual}} - A_{\text{dual}}$ has spectral gap $\lambda_2 = 2.138$ (vs $\mathbb{Z}^3$ ’s $\lambda_2 = 3.0$ on the same number of vertices), and the spectral support extends to $\lambda_{\text{max}} = 18.7$ (vs $\mathbb{Z}^3$ ’s 12). The dual graph is denser (60% more edges per vertex) but more weakly connected at low modes (29% smaller spectral gap) than the cubic dual lattice. As a sanity check, the $3 \times 3 \times 3$ cubic dual graph reproduces avg degree = 6.00, $\lambda_2 = 3.00$ (analytic DFT value), and the 27 monopole sites have cube-boundary closure verified to machine precision; the standard BMK derivation is recovered exactly when Rule B is replaced by $\mathbb{Z}^3$.

Where the dual Laplacian enters the dilute-gas formula. The standard Polyakov derivation [13, 14] is a sine-Gordon dualization of monopole-instanton configurations. The dual scalar field χ lives on dual lattice sites; the kinetic term is $(\nabla\chi)^\top \nabla\chi$ with ∇ the dual-lattice gradient. Quadratic expansion of the cosine-interaction term around $\chi = 0$ gives a Debye mass $m_D^2 = 4\pi^2\rho_M/\beta_V$, and the resulting screened-Coulomb propagator gives the Polyakov string-tension formula $\sigma(\beta) = (4/\pi)\sqrt{\rho_M/\beta_V}$.

The non-cubic dual graph alters the kinetic term: $(\nabla\chi)^\top \nabla\chi$ is now generated by L_{dual} rather than the $\mathbb{Z}^3$ Laplacian. But the kinetic term and the cosine interaction live on the dual lattice *together*, and the Debye mass m_D entering the string tension is the mass eigenvalue of the quadratic-around-vacuum theory; it depends on the *ratio* of the cosine coupling to the kinetic coefficient, not on the absolute spectral gap. The β -dependence of m_D^2 is therefore still $m_D^2 \propto \rho_M$, and the

β -dependence of σ is still $\sigma \propto \sqrt{\rho_M}$. **The rate constant $c_\sigma = c_\rho/2$ is invariant under the dual-lattice topology change; only the absolute scale σ_0 inherits the dual-lattice structure.** Specifically, the Rule B prefactor scales with $1/\sqrt{\lambda_2^{\text{dual}}}$, which gives $\sigma_0^{\text{RuleB}}/\sigma_0^{\mathbb{Z}^3} \approx \sqrt{3.0/2.138} \approx 1.18$, an $O(1)$ correction to the prefactor that does not affect the rate.

What XCWG-G actually measured. The numerical identity $c_{\sigma_{\text{ens}}} \approx c_{\mu_{\text{comb}}}$ resolves the question. Direct computation on the XCWG-G $n_{\text{max}} = 2$ MC data gives $c_{\sigma_{\text{ens}}} = 1.196$ and $c_{\mu_{\text{comb}}} = 1.242$, agreement to 4% – well within MC noise. The two coefficients move together because at $n_{\text{max}} = 2$ the available Wilson loops have $L \approx 4A$ for most (A, L) pairs in the ensemble (cluster geometry forces approximate proportionality), so the perimeter sector and the fit-extracted ensemble-mean “area sector” track each other. **The “ $c_\sigma = 1.20$ ” measured by XCWG-G is the perimeter-coefficient rate, not the string-tension rate.** The true area-coefficient σ_{comb} in the joint fit is statistically zero at every β ($|\sigma_{\text{comb}}| < 0.011$ across $\beta \in [0.1, 10]$, mean -0.0036), and this is fully consistent with the Polyakov-predicted $\sigma(\beta = 1.5) \approx 5 \times 10^{-9}$ sitting four orders below the MC detection floor. No disagreement exists when the observables are correctly identified.

Finite-volume floor. The dual graph at $n_{\text{max}} = 2$ has finite volume $V_{\text{dual}} = 60$, which gives a soft string-tension floor at $\sigma_{\text{floor}} \approx 0.658$ from the IR-truncated Coulomb propagator (the would-be infrared-divergent integral becomes finite on a finite graph). The crossover scale is $\beta \approx 0.28$. Above this β , the Polyakov-predicted σ is sub-floor; below it, the prediction approaches the floor smoothly. The XCWG-G observation $\sigma_{\text{ens}}(\beta = 0.3) \approx 0.17$ – well below the floor – is the perimeter sector again, not the floor-limited Polyakov prediction.

Verdict. The v4 caveat (vi) was a misidentification. The disagreement between $c_\rho = 9.40$ (XCWG-F) and the apparent $c_\sigma^{\text{MC}} = 1.20$ (XCWG-G) is not a missing structural correction; it is two different observables. The Polyakov rate $c_\sigma = c_\rho/2 = 4.70$ is correct on Rule B; the corresponding σ_{Polyakov} is unobservably small at the available MC statistics, and $\sigma_{\text{comb}} \approx 0$ is exactly what is expected. Non-cubic dual-lattice corrections are real but enter the prefactor σ_0 via the dual Laplacian spectral gap, not the rate. The framework is internally self-consistent; the v5 update converts caveat (vi) from “open quantitative gap” to “structurally resolved.”

7 Seven-witness structural verdict

The seven independent diagnostics give:

Witness	Diagnostic	Verdict
1 (§4)	Heat-kernel d_s at $n_{\max} = 3, 4, 5$ Weyl d_W at $n_{\max} = 3, 4, 5$ Scalar Fock d_s comparison	$1.86 \rightarrow 2.27 \rightarrow 2.54$ (toward 3) $2.48 \rightarrow 3.46 \rightarrow 3.72$ (crosses 3 at $n_{\max}$) $1.68 \rightarrow 1.76 \rightarrow 1.79$ (plateaus near 2)
2 (§5)	MK β_{eff} flow at $n_{\max} = 2, 3$ Asymptotic-limit verification	monotone to $\beta = 0$, no UV fixed point match strong/weak coupling within $\leq$
3 (§6.1)	Gaussian $\langle S(L) \rangle \sim L^\alpha$ at $n_{\max} = 2, 3, 4, 5$	$\alpha \in [0.89, 1.18]$ = 3D free-photon perimeter law
4 (§6.2)	Strong-coupling area law $\langle W \rangle = (I_1/I_0)^A$ Asymptotic-limit verification Area-vs-perimeter discriminator at $n_{\max} = 2, 3, 4$	$\sigma(\beta) > 0$ at every $\beta > 0$ $\sigma \rightarrow \ln(2/\beta)$, $1/(2\beta)$ to 5×10^{-4} $\langle W \rangle$ depends only on A , not on L (at
5 (§6.4)	Monopole $\rho_M(\beta)$ via DeGrand–Toussaint at $n_{\max} = 2, 3$ Polyakov dilute-gas fit $\rho_M = A e^{-c\beta}$ Size-scaling consistency Floor-crossover check 3D dilute-plasma vs 4D condensation	exponential decay over 5 decades $c = 9.40$ ($n = 2$), 8.95 ($n = 3$), $R^2 = 0.99$ c $n_{\max}$ -independent within 5% $\beta_{\text{floor}}^{\text{pred}} \approx 1.1$ matches obs. $\beta_{\text{zero}} = 1.0$ – 1.1 no zero crossing, monotone expo
6 (§6.5)	Full-MC Wilson loops at $n_{\max} = 2, 3$ $\sigma_{\text{ens}}(\beta)$ at $\beta \in [0.1, 10]$ Joint area+perimeter fit $\log W = -\sigma A - \mu L + c$ Polyakov-rate cross-check Verdict	$\sigma_{\text{ens}}(\beta) > 0$ monotone decreasing $0.377 \rightarrow 0.011 \rightarrow 0.003$ (monotone) $\sigma_{\text{comb}} \approx 0$, $\mu_{\text{comb}} \approx \sigma_{\text{ens}}$ exp. shape $R^2 = 0.83$ – 0.88 , rate 3 – $4 \times$ WEAK PASS, perimeter-dominant
7 (§6.6)	Polyakov loop on $G_B \times C_{N_t}$ at $n_{\max} = 2$, $N_t \in \{4, 8\}$ Homological Polyakov class $N_t = 2$ kinematic degeneracy Two-point correlator decay 3D persistent confinement vs 4D Svetitsky-Yaffe	$ \langle P \rangle = 0$ within 2σ at all 6 cells $\dim \ker(d_1) = 1$ at $N_t \geq 3$ (forced cocycle) $P \equiv 1$ identically, $\dim \ker(d_1) = 0$ (exact) exponential in d , per-step ratio 0.72 at $d = 4$ no deconfinement transition, hot

Joint structural verdict (promoted from five-witness to seven-witness). The seven witnesses each return values consistent with 3D compact $U(1)$ (Polyakov–BMK universality class with no finite-temperature deconfinement transition) and inconsistent with 4D compact $U(1)$. The verdicts are independent: spectral dimension is a property of L_0 alone; MK flow is a property of the character expansion under bond-moving; Gaussian Wilson loop probes the co-exact Hodge sector $K = d_1^\top d_1$; strong-coupling Wilson loop probes the full character expansion at leading order; the monopole density probes the non-perturbative compact integration domain through the canonical Polyakov discriminator; the full-MC Wilson loops probe the same compact action at all orders in the character expansion simultaneously; the Polyakov loop on $G_B \times C_{N_t}$ probes the finite-temperature behaviour via the homologically forced winding observable. *Witnesses 1–4 verify the leading-order signatures shared by 3D and 4D compact $U(1)$; witnesses 5–7 distinguish them along three structurally distinct non-perturbative axes (dilute-monopole gas, full compact integration, and Kaluza-Klein reduction), and Rule B falls on the 3D side in each.*

Six clean passes plus one nuanced weak-pass. Witnesses 1, 2, 3, 4, 5, and 7 are clean passes at the level each test accesses. Witness 6 is a nuanced weak-pass: the naive ensemble-mean string tension $\sigma_{\text{ens}}(\beta)$ is positive and monotone decreasing at every β , qualitatively consistent with 3D

compact $U(1)$, but a joint area-plus-perimeter fit on per-loop data reveals that the apparent confinement signal is carried by the perimeter coefficient μ_{comb} rather than the area coefficient σ_{comb} on the available graph sizes. The v4 expectation was that this would shrink at $n_{\text{max}} \geq 4$; the v5 sprint XCWG-G $n_{\text{max}} = 4$ extension sharpens that diagnosis (see §6.5): the negative σ_{comb} moves *deeper* with n_{max} , not weaker, because the area-controlled-cluster construction itself produces a non-monotonic $A \leftrightarrow L_{\text{min}}$ coupling that no linear joint fit can disentangle. The honest area-law coefficient at fixed L is consistent with zero at the available MC statistics, consistent with the Polyakov-predicted $\sigma \sim 5 \times 10^{-9}$ sitting four orders below the MC detection floor (XCWG-J, §6.7). The structurally-compatible reading survives because (a) all six other witnesses pass cleanly, (b) the weak-pass deviation is along a cluster-construction methodological axis rather than indicating 4D-like physics, and (c) the Polyakov rate constant agreement is structurally exact at the rate level once the observable identification is corrected – it is the prefactor σ_0 that carries non-cubic dual-lattice corrections, not the rate.

The five-to-seven promotion converts the v3 reading “structurally compatible with 3D compact $U(1)$ at leading order and non-perturbatively via the Polyakov monopole-density signature” into the v4 reading: **structurally compatible with 3D compact $U(1)$ at leading order, non-perturbatively via the Polyakov monopole-density signature, full-MC Wilson-loop qualitative confinement signature (with finite-volume scope on the area-vs-perimeter separation), and persistent confinement at finite temperature via the homological Polyakov-class structure.** This is the strongest cumulative case for 3D-compatibility achievable from a sprint-scale investigation at $n_{\text{max}} \in \{2, 3\}$. It does not collapse to “Rule B is 3D compact $U(1)$ ” because the non-perturbative witnesses operate at sprint-level resolution rather than the thermodynamic limit, and witness 6’s clean area-vs-perimeter separation is deferred to $n_{\text{max}} \geq 4$; what is established is that the canonical 3D-compact- $U(1)$ structural signatures are present, not that all higher-order or strong-correlation features of the 3D universality class are reproduced.

What this verdict does *not* establish.

- (i) *Unique identification with 3D compact $U(1)$ up to universality.* Witnesses 1–4 are leading-order or coarse-approximation tests; witnesses 5–7 are non-perturbative but at sprint-level Monte Carlo resolution ($n_{\text{max}} \in \{2, 3\}$, β -grid in 10–15 points, sample sizes $\sim 10^3$). The thermodynamic limit is not directly probed; what is probed is internal consistency between two cutoffs and the categorical absence of 4D-like signatures (finite- β_c zero crossing in ρ_M ; deconfinement transition in $\langle P \rangle$). These together exclude the 4D phase-transition scenario at the sprint’s resolution but do not exclude all other universality classes that share the 3D dilute-plasma signature at the order tested.
- (ii) *Continuum 3D compact QED.* Per CLAUDE.md §1.5, the Rule B Wilson construction is a finite-graph abelian lattice gauge theory. The continuum limit $n_{\text{max}} \rightarrow \infty$ has not been computed; whether it exists and reproduces 3D compact QED on a Riemannian manifold is open.
- (iii) *A continuum field-theoretic interpretation of Rule B.* The construction is a structural discrete analog. The graph and the continuum are dual descriptions connected by Paper 7’s Fock projection; this paper does not assert ontological priority of either.
- (iv) *Quantitative matching of the Polyakov constant v_0 to the cubic-lattice baseline.* The XCWG-F fit returns $v_0^{\text{Rule B}} \approx 0.466$, of order unity but not equal to the cubic-lattice value; the irregular non-cubic geometry of Rule B makes a precise numerical match unexpected, and no

such match is claimed. The structural-compatibility statement is that the *form* of the decay is the Polyakov 3D form, not that the numerical v_0 should agree with $\mathbb{Z}^3$.

- (v) *Cluster-topology-induced $A \leftrightarrow L$ coupling in the full-MC Wilson loop ensemble.* Witness 6 (full MC at $n_{\max} \leq 4$) gives $\sigma_{\text{ens}}(\beta) > 0$ monotone decreasing – qualitatively the right 3D compact $U(1)$ shape – but the joint $\sigma A + \mu L$ fit returns σ_{comb} statistically negative at high significance, sharpening *with* $n_{\max}$ rather than fading ($z = -2.0, -4.8, -48$ at $n_{\max} = 2, 3, 4$). The v5 sharpening (§6.5) traces this to the area-controlled-cluster construction: at $n_{\max} = 4$ the $L_{\min}(A)$ sequence is 4, 4, 6, 8, 12, 10 for $A = 1 \dots 6$, which is *non-monotonic*, producing the 20σ effect $\langle W(A=4) \rangle > \langle W(A=3) \rangle$ at $\beta = 0.3$ – anti-area-law in the ensemble, but driven by cluster topology, not gauge dynamics. The linear $\sigma A + \mu L$ ansatz cannot disentangle this coupling. The clean fixed- $L = 6$ slice at $\beta = 0.3$ gives $\sigma_{\text{fix-}L=6} = -0.004 \pm 0.003$, consistent with zero – the honest area-law coefficient at the available MC statistics, and fully consistent with the Polyakov-dilute-gas prediction $\sigma_{\text{Polyakov}} \sim 5 \times 10^{-9}$ at the same β (four orders below the MC detection floor). The diagnosis is therefore not finite-volume “larger graph will fix it” but a structural property of the cluster construction on Rule B’s dense-plaquette graph: the next sharpening requires either extended MC statistics (push the floor below the Polyakov prediction) or fixed- L ensemble construction by design. The qualitative 3D-compatibility reading is unaffected; the strong area-law verdict is carried by witnesses 4 (LO character expansion) and 5 (monopole density), which use methodologies immune to the cluster-topology issue.
- (vi) *Polyakov rate-constant agreement – RESOLVED in v5.* The v4 caveat recorded a 3–4 $\times$ disagreement between XCWG-F’s $c_\rho = 9.40$ and witness 6’s $c_\sigma^{\text{MC}} = 1.20$, attributed tentatively to non-cubic dual-lattice topology. Sprint XCWG-J (§6.7) refines the question and resolves it: the disagreement is a category mismatch between observables, not a missing structural correction. Direct numerical computation on XCWG-G’s $n_{\max} = 2$ MC data gives $c_{\sigma_{\text{ens}}} = 1.196$ and $c_{\mu_{\text{comb}}} = 1.242$, agreement to 4%, so what witness 6 measures as the “string tension rate” is in fact the perimeter-coefficient rate. The true Polyakov-predicted area-coefficient σ_{comb} is statistically zero at every β ($|\sigma_{\text{comb}}| < 0.011$, mean -0.0036), which is exactly the prediction $\sigma_{\text{Polyakov}}(\beta = 1.5) \sim 5 \times 10^{-9}$ sitting four orders below the MC detection floor. Non-cubic dual-lattice corrections do exist (dual graph average degree 9.60 vs $\mathbb{Z}^3$ ’s 6.0, dual Laplacian spectral gap $\lambda_2 = 2.138$ vs 3.0), but in the standard sine-Gordon dualization they enter the prefactor σ_0 , not the rate constant $c_\sigma = c_\rho/2$. The framework is internally self-consistent on the Polyakov rate; the structural-compatibility statement is strengthened, not weakened, by the resolution.

The combination of (i)–(vi) keeps the structural-compatibility framing as the load-bearing reading: *Rule B is structurally compatible with 3D compact $U(1)$ on a non-cubic graph at leading order, non-perturbatively via the Polyakov dilute-monopole-plasma signature, qualitatively at full Monte Carlo on the compact Wilson measure (subject to finite-volume perimeter-dominance), and at finite temperature via the homologically- forced Polyakov-loop persistent-confinement signature*, as established by seven independent witnesses at the order each test accesses (six clean passes, one nuanced weak-pass). The five-to-seven promotion is a genuine strengthening of the v3 verdict – two structurally distinct non-perturbative observables flagged as outlook in v3 (full-MC Wilson loops and Polyakov loops on $G_B \times C_{N_t}$) have been executed, with one clean pass and one nuanced weak-pass – without crossing into a unique-identification claim that the data do not support.

8 Outlook

Several directions sharpen the structural-compatibility verdict toward unique identification or refine the open questions. Six completed sprints (XCWG-E, XCWG-F, XCWG-G, XCWG-H, plus the v5-update sprints XCWG-G at $n_{\max} = 4$ and XCWG-J Polyakov-rate refinement) are summarised first. The v5 sprints closed the two v4 honest-scope caveats structurally: caveat (v) sharpens to a cluster-topology-induced $A \leftrightarrow L$ coupling (not finite-volume), and caveat (vi) resolves as a category mismatch between observables, not a missing structural correction. With both v4 caveats now in structurally-resolved status, the cleanest next single-agent sprint is either (a) extended MC statistics to push the detection floor below the Polyakov-predicted $\sigma \sim 5 \times 10^{-9}$ (would convert $\sigma_{\text{comb}} = 0$ from a consistency statement to a direct measurement), or (b) a fixed- L ensemble construction by design rather than cluster-area control, which would isolate the area-coefficient from cluster geometry.

XCWG-E NLO character expansion: completed with mixed verdict (supplied combinatorics for XCWG-F). The next-to-leading-order character-expansion follow-up was executed in May 2026 and is reported in §6.3. It produced one substantive new structural finding – the smallest closed 2-cycle on Rule B has size $k_\delta = 3$, smaller than $\mathbb{Z}^4$ ’s $k_\delta = 6$, reflecting the graph’s high plaquette density – which directly supplied the elementary monopole-site combinatorics that XCWG-F’s DeGrand–Toussaint construction needed. The direct NLO Wilson-loop diagnostic was truncation-dominated, but the Creutz-like $\sigma_{\text{eff}}(2 \rightarrow 3)$ stayed positive at all β (consistent with the LO 3D-like reading), and the $k_\delta = 3$ finding became load-bearing for the fifth-witness construction.

XCWG-F monopole density: completed with clean pass (§6.4). Polyakov’s mechanism for permanent confinement in 3D compact $U(1)$ goes through the magnetic-monopole plasma: in $D = 3$ the monopole density $\rho_M(\beta)$ stays nonzero at all β as a dilute plasma; in $D = 4$ the same density drops sharply to zero above a critical coupling $\beta_c \approx 1.01$ as the monopoles condense into a dilute gas that admits the deconfined phase [15, 13]. The DeGrand–Toussaint construction [24] adapted to Rule B’s size-3 triangle-prism elementary sites, combined with a Metropolis–Hastings Monte Carlo of compact $U(1)$ Wilson, gives $\rho_M(\beta)$ decaying exponentially over five decades with Polyakov fit $\rho_M = A e^{-c\beta}$ at $R^2 = 0.981$ and $n_{\max}$ -independent decay rate $c \approx 9.2$ within 5%. The fifth witness passes; the v2 verdict “deferred non-perturbative discriminator” is now closed.

XCWG-G full Monte Carlo Wilson loops: completed with nuanced weak-pass and a substantive finite-volume finding (§6.5). Full Metropolis–Hastings MC of the compact $U(1)$ Wilson measure on Rule B at $n_{\max} \in \{2, 3\}$ returns $\sigma_{\text{ens}}(\beta) > 0$ monotone decreasing at every $\beta \in [0.1, 10]$, qualitatively consistent with 3D compact $U(1)$. But the joint $\log W = -\sigma A - \mu L + c$ fit on per-loop data reveals that the apparent confinement signal is carried by the perimeter coefficient μ_{comb} , not the area coefficient σ_{comb} : on the available graph sizes, full-MC Wilson loops are perimeter-dominated. The Polyakov cross-check $\sigma \propto \exp(-(c_\rho/2)\beta)$ from XCWG-F passes qualitatively (exponential shape, $R^2 = 0.83\text{--}0.88$) but quantitatively deviates by a factor of 3–4 $\times$ in the rate constant, attributable to Rule B’s non-cubic dual-lattice topology preventing the textbook continuum Polyakov derivation from applying verbatim. The sixth witness passes weakly with finite-volume scope; a clean area-vs-perimeter separation requires $n_{\max} \geq 4$. This is the cleanest next-sprint target (see below).

XCWG-H Polyakov loops on $G_B \times C_{N_t}$: completed with clean pass and a structurally novel homological mechanism (§6.6). A brand-new product-graph construction $G_B \times C_{N_t}$ at $N_t \in \{4, 8\}$ has $\dim \ker(d_1) = 1$, with the one extra cocycle being exactly the Polyakov-loop class – the homological signature that makes $\langle P \rangle$ a well-defined observable not by analogy to continuum 4D $U(1)$ but as a homologically forced cocycle of the discrete object. The $N_t = 2$ case is a kinematic degeneracy ($P \equiv 1$ identically by single-edge backtracking; $\dim \ker(d_1) = 0$), correctly excluded from the verdict. MC at six non-degenerate (β, N_t) cells returns $|\langle P \rangle| = 0$ within 2σ at every cell; the two-point Polyakov correlator decays exponentially with spatial graph distance, with the correct N_t -dependence (hotter systems screen further). The seventh witness passes cleanly: Rule B Wilson exhibits persistent confinement at finite temperature, exactly the 3D compact $U(1)$ signature that distinguishes it from the 4D Svetitsky–Yaffe deconfinement transition.

XCWG-G $n_{\max} = 4$ extension: completed in v5 with sharpened diagnosis (§6.5, $n_{\max} = 4$ paragraph). Full MC on Rule B at $n_{\max} = 4$ ($V = 60$, $E = 312$, $\beta_1 = 253$, 30,600 sweeps per β) was expected to clean up the joint $\sigma A + \mu L$ fit’s $\sigma_{\text{comb}} \approx 0$. The actual finding is more structural: σ_{comb} moves *deeper* negative at $z = -48$ (vs $-2.0, -4.8$ at $n_{\max} = 2, 3$), driven by a non-monotonic $L_{\min}(A) = 4, 4, 6, 8, 12, 10$ for $A = 1 \dots 6$ in the area-controlled cluster construction, producing the 20σ effect $\langle W(A = 4) \rangle > \langle W(A = 3) \rangle$ at $\beta = 0.3$. The clean fixed- $L = 6$ slice at $\beta = 0.3$ gives $\sigma = -0.004 \pm 0.003$ (consistent with zero), which is fully consistent with the Polyakov-predicted $\sigma \sim 5 \times 10^{-9}$ at the same β – four orders below the MC detection floor. Diagnosis sharpens from “finite-volume, will fade” to “cluster-topology-induced non-linear $A \leftrightarrow L$ coupling that the joint linear fit cannot resolve at any tested graph size; honest area-law signal is at or below MC noise.” The next sharpening requires either (a) extended MC statistics to push the detection floor below the Polyakov-predicted σ , or (b) a fixed- L ensemble construction by design rather than cluster-area control.

XCWG-J Polyakov-rate refinement: completed in v5 with structural resolution (§6.7). The v4 caveat (vi) was the apparent 3–4 $\times$ disagreement between $c_\rho = 9.40$ (XCWG-F) and $c_\sigma^{\text{MC}} = 1.20$ (XCWG-G). The XCWG-J refinement built the Rule B dual graph G_B^* at $n_{\max} = 2$ (average degree 9.60, dual Laplacian spectral gap $\lambda_2 = 2.138$; cubic-baseline BMK reproduced exactly on the $3 \times 3 \times 3$ sanity check). The verdict: the disagreement is not a missing structural correction – it is a category mismatch between two different observables. Direct computation of $c_{\sigma_{\text{ens}}} = 1.196$ vs $c_{\mu_{\text{comb}}} = 1.242$ (4% agreement) shows XCWG-G measures the perimeter-coefficient rate, not the Polyakov-predicted string-tension rate. The latter $\sigma_{\text{Polyakov}}(\beta \geq 1.5) \sim 5 \times 10^{-9}$ is below the MC detection floor, consistent with the $\sigma_{\text{comb}} \approx 0$ measurement. Non-cubic dual-lattice corrections are real but enter the prefactor σ_0 via the dual Laplacian, not the rate $c_\sigma = c_\rho/2$. The framework is internally self-consistent on the rate; caveat (vi) closes as resolved.

Other sharpenings of completed sprints. Three further follow-ons of more modest scope: (a) longer XCWG-F MC runs at large β to push the detection floor below 10^{-7} , directly measuring ρ_M in the regime currently inferred from exponential extrapolation; (b) analytical computation of the XCWG-F dual cell-volume v_0 on Rule B from the plaquette-adjacency structure, enabling a sharper $v_0^{\text{Rule B}}$ vs cubic-lattice baseline comparison than the current order-of-magnitude statement $v_0^{\text{Rule B}} \approx 0.466$; (c) XCWG-H extension to $n_{\max} = 3$ at multiple N_t to extract a clean two-point Polyakov correlator decay rate (the $n_{\max} = 2$ graph diameter is only 4, which caps the correlator’s dynamic range). Each is ~ 1 –3 weeks of focused effort.

Beyond the $U(1)$ Wilson program. With the seven-witness verdict now solid, the natural next direction is to ask whether non-abelian Wilson gauge theory ($SU(2)/SU(3)$) admits a Rule B-analog construction. Paper 30’s $SU(2)$ Wilson lives on the scalar Hopf S^3 graph (per- ℓ 2D, same structural limitation as Paper 25); Sprint ST- $SU(3)$ ’s $SU(3)$ Wilson lives on the Bargmann–Segal S^5 graph and is pure-gauge (no natural matter coupling). The cross- ℓ Dirac structure of Rule B suggests a non-abelian generalization: replacing the $U(1)$ link variables with $SU(N)$ -valued links would give a Yang–Mills-on-Rule-B construction whose 3D phase content could be tested against the same seven-witness diagnostic suite. This is multi-month rather than single-sprint scope but is well-defined in principle.

Position in the GeoVac Wilson program. Combining this paper with the existing constructions:

Paper	Gauge group	Substrate	Cross- ℓ ?
25 [5]	$U(1)$	scalar Hopf S^3 graph	no (per- ℓ 2D)
30 [8]	$SU(2)$	scalar $S^3 = SU(2)$ graph	no
Sprint ST- $SU(3)$	$SU(3)$	Bargmann–Segal S^5 graph	—
41 (this paper)	$U(1)$	Dirac Rule B graph	yes (E1 dipole)

The four constructions instantiate all SM compact gauge groups ($U(1) \times SU(2) \times SU(3)$) as Wilson lattice gauge theories on natural GeoVac sub-manifolds, with Rule B supplying the substrate on which the abelian theory has its 3D phase content.

Relation to Papers 38, 40. The propinquity convergence theorems of Papers 38 [10] and 40 [11] establish that finite GeoVac graph triples converge to the round- S^3 spectral triple in the Latrémolière propinquity, with universal rate constant $4/\pi$. The present paper’s Rule B graph is a Dirac-spinor lift of the scalar Fock graph on which Papers 38 and 40 work; a propinquity analysis of the Rule B Wilson construction would parallel those papers’ arguments but is out of scope here. The relevant cross-reference: Paper 40’s universality of the rate constant suggests that the asymptotic geometry of Rule B and the scalar Fock graph agree in the continuum limit, which would be consistent with the spectral-dimension finding that Rule B trends toward $d_s = 3$ while the scalar graph plateaus at $d_s \approx 2$ – the latter being a finite- $n_{\max}$ property that washes out in the continuum.

Relation to Marcolli–van Suijlekom gauge networks. The present construction sits inside the gauge-networks-in-NCG framework of Marcolli and van Suijlekom [27], with the continuum limit clarified by Perez-Sanchez [25, 26] as Yang–Mills without an autonomously generated Higgs sector (the Higgs is supplied by an external off-diagonal D_F choice, not by GeoVac substrate data; cf. Paper 32 § VIII.C Sprint H1 verdict). The Rule B Wilson $U(1)$ construction reported here is a discrete-substrate instance of this lineage on the Dirac-spinor lift of the Fock graph; the abelian-vs-non-abelian split is the maximal-torus projection examined in Paper 30 [8].

Acknowledgements

The construction extends Paper 25’s $U(1)$ Wilson on the scalar Hopf graph and Paper 29 §RH-C’s Dirac Rule B definition. The four-witness sprint sequence (XCWG-A foundation; XCWG-B RG + predictions + observables; XCWG-C Gaussian Wilson loop; XCWG-D strong-coupling

Wilson loop) was executed in May 2026; the v2 update reported the XCWG-E NLO character-expansion follow-up of the same month, adding the $k_\delta = 3$ structural finding and the Creutz-ratio cross-check (§6.3); the v3 update reported the XCWG-F non-perturbative monopole-density sprint, using DeGrand–Toussaint extraction on Rule B’s size-3 triangle-prism elementary sites and Metropolis–Hastings Monte Carlo of compact $U(1)$ Wilson (§6.4). The v4 update reports two further sprints executed in May 2026: XCWG-G (full Monte Carlo Wilson loops on the compact $U(1)$ measure, §6.5), which delivers a nuanced weak-pass sixth witness and identifies finite-volume perimeter-dominance on Rule B at $n_{\max} \leq 3$ as the load-bearing follow-on; and XCWG-H (Polyakov loops on a brand-new product-graph construction $G_B \times C_{N_t}$, §6.6), which delivers a clean seventh-witness pass for persistent confinement at finite temperature via the homological Polyakov-loop class. All numerical data are reproducible from the driver scripts and JSON output files in the `debug/` directory of the GeoVac repository, including `debug/xcwg_nlo_character_expansion.py` and `debug/data/xcwg_nlo_character_expansion.json` for the v2 NLO results; `debug/xcwg_monopole_density.py` together with `debug/data/xcwg_monopole_density.json` and `debug/plots/xcwg_monopole_density.png` for the v3 monopole-density results; `debug/xcwg_full_mc_wilson_loops.py` together with `debug/data/xcwg_full_mc_wilson_loops.json` and `debug/plots/xcwg_full_mc_wilson_loops.png` for the v4 XCWG-G full-MC Wilson-loop results; and `debug/xcwg_polyakov_loop.py` together with `debug/data/xcwg_polyakov_loop.json` for the v4 XCWG-H Polyakov-loop results. Per CLAUDE.md §1.5, no ontological priority is asserted for the discrete graph over the continuum manifold; the two are dual descriptions connected by Paper 7’s Fock projection.

References

- [1] GeoVac Project, “The Dimensionless Vacuum: S^3 Conformal Equivalence and 18 Symbolic Proofs,” Paper 7 (2025).
- [2] GeoVac Project, “Structurally Sparse Qubit Hamiltonians from Natural Geometry,” Paper 14 (2025).
- [3] GeoVac Project, “Spectral-Geometric Exchange Constants,” Paper 18 (2026).
- [4] GeoVac Project, “The Bargmann–Segal Lattice: π -Free Discretization of the 3D Harmonic Oscillator on S^5 ,” Paper 24 (2026).
- [5] GeoVac Project, “The Hopf Graph as a Lattice Gauge Structure: A Framework Observation,” Paper 25 (2026).
- [6] GeoVac Project, “QED on S^3 : Spectral Geometry of Perturbative Transcendentals,” Paper 28 (2026).
- [7] GeoVac Project, “The GeoVac Hopf Graphs Are Ramanujan: Ihara Zeta, Graph-RH, and a Scope Boundary on Selberg-on-Hydrogen,” Paper 29 (2026).
- [8] GeoVac Project, “SU(2) Wilson Lattice Gauge Structure on the GeoVac Hopf Graph: The Non-Abelian Sibling of Paper 25,” Paper 30 (2026).
- [9] GeoVac Project, “Time as Projection: Klein–Gordon, Matsubara, and the Master Mellin Engine,” Paper 35 (2026).
- [10] GeoVac Project, “SU(2) Propinquity Convergence on the Camporesi–Higuchi Spectral Triple,” Paper 38 (2026).

- [11] GeoVac Project, “Unified GH-Convergence on Compact Lie Groups: Universal Rate Constant $4/\pi$ via Plancherel $\times$ Vandermonde Cancellation,” Paper 40 (2026).
- [12] K. G. Wilson, “Confinement of quarks,” *Phys. Rev. D* **10**, 2445–2459 (1974).
- [13] A. M. Polyakov, “Quark confinement and topology of gauge theories,” *Nucl. Phys. B* **120**, 429–458 (1977).
- [14] T. Banks, R. Myerson, and J. Kogut, “Phase transitions in Abelian lattice gauge theories,” *Nucl. Phys. B* **129**, 493–510 (1977).
- [15] A. H. Guth, “Existence proof of a nonconfining phase in four-dimensional $U(1)$ lattice gauge theory,” *Phys. Rev. D* **21**, 2291–2307 (1980).
- [16] J. Fröhlich and T. Spencer, “The Kosterlitz–Thouless transition in two-dimensional Abelian spin systems and the Coulomb gas,” *Commun. Math. Phys.* **81**, 527–602 (1981); *Phys. Rev. Lett.* **46**, 1006–1009 (1981).
- [17] A. A. Migdal, “Phase transitions in gauge and spin-lattice systems,” *Sov. Phys. JETP* **42**, 743 (1975).
- [18] L. P. Kadanoff, “Notes on Migdal’s recursion formulas,” *Ann. Phys.* **100**, 359–394 (1976).
- [19] M. Creutz, *Quarks, Gluons and Lattices*, Cambridge University Press, Cambridge (1983).
- [20] M. N. Chernodub, E.-M. Ilgenfritz, and A. Schiller, “A lattice study of 3D compact QED at finite temperature,” arXiv:hep-lat/0105021 (2001).
- [21] M. Loan and C. Hamer, “Static quark potential and string tension for compact $U(1)$ in $(2+1)$ dimensions,” arXiv:hep-lat/0208047 (2002).
- [22] L.-H. Lim, “Hodge Laplacians on graphs,” *SIAM Review* **62**, 685–715 (2020).
- [23] M. T. Schaub, A. R. Benson, P. Horn, G. Lippner, and A. Jadbabaie, “Random walks on simplicial complexes and the normalized Hodge 1-Laplacian,” *SIAM Review* **62**, 353–391 (2020).
- [24] T. A. DeGrand and D. Toussaint, “Topological excitations and Monte Carlo simulation of Abelian gauge theory,” *Phys. Rev. D* **22**, 2478–2489 (1980).
- [25] C. I. Perez-Sanchez, “Bratteli networks and the Spectral Action on quivers,” arXiv:2401.03705 (2024).
- [26] C. I. Perez-Sanchez, “Comment on ‘Gauge networks in noncommutative geometry’,” arXiv:2508.17338 (2025).
- [27] M. Marcolli and W. D. van Suijlekom, “Gauge networks in noncommutative geometry,” *J. Geom. Phys.* **75**, 71–91 (2014), arXiv:1301.3480.